

Elaine: Elaborations of Higher-Order Effects as First-Class Language Feature



Terts Diepraam
September 11, 2023

Elaine: Elaborations of Higher-Order Effects as First-Class Language Feature

THESIS

submitted in partial fulfillment of the
requirements for the degree of

MASTER OF SCIENCE

in

COMPUTER SCIENCE

by

Terts Diepraam
born in Amsterdam, the Netherlands



Programming Languages Group
Department of Software Technology
Faculty EEMCS, Delft University of Technology
Delft, the Netherlands
www.ewi.tudelft.nl

© 2023 Terts Diepraam.

Cover picture: Rubin's Vase.

Elaine: Elaborations of Higher-Order Effects as First-Class Language Feature

Author: Terts Diepraam
Student id: 5652235
Email: t.diepraam@student.tudelft.nl

Abstract

Thesis Committee:

Chair: Prof. dr. M. Spaan, Faculty EEMCS, TU Delft
Committee Member: Dr. C. Bach Poulsen, Faculty EEMCS, TU Delft
University Supervisor: Ir. C. van der Rest, Faculty EEMCS, TU Delft

Contents

Contents	iii
1 Introduction	1
1.1 Contributions	3
1.2 Artefact	3
2 Algebraic Effects	5
2.1 Monads	5
2.2 Effect Composition with the Free Monad	8
2.3 Algebraic Effects	11
2.4 Building a Language with Algebraic Effects	12
3 Higher-Order Effects	15
3.1 Computation Parameters	15
3.2 Breaking Algebraicity	16
3.3 The Modularity Problem	17
3.4 Hefty Algebras	17
4 A Tour of Elaine	21
4.1 Basics	21
4.2 Types	22
4.3 Algebraic Data Types	24
4.4 Recursion, Loops and Lists	25
4.5 Algebraic Effects	26
4.6 Effect-Agnostic Functions	30
4.7 Higher-Order Effects	31
5 Implicit Elaboration Resolution	35
6 Elaboration Compilation	37
6.1 Non-locality of Elaborations	37
6.2 Operations as Functions	38
6.3 Compiling Elaborations to Dictionary Passing	38
6.4 Compiling Elaborations into Handlers	39
7 Related Work	43
7.1 Monad Transformers	43
7.2 Effects as Free Monads	44
7.3 Prior Art for First-Class Effects	45

8 Conclusion	47
Bibliography	49
A Elaine Example Programs	53
A.1 A naive SAT solver	53
A.2 The Reader Effect	54
B Elaine Specification	57
B.1 Syntax definition	57
B.2 Effect row semantics	57
B.3 Typing judgments	60
B.4 Desugaring	63
B.5 Semantics	63
B.6 Standard Library	64

Chapter 1

Introduction

TODO: Start: effects are a thing

TODO: Examples in Koka throughout!

In many programming languages, computations are allowed to have *effects*. This means that they can do things besides producing output and interact with their environment. It might, for instance, read or modify a global variable, write to a file, throw an exception or even exit the program altogether. These are all examples of effects.

Historically, programming languages have supported effects in different ways. Some programming languages opt to give the programmer virtually unrestricted access to effectful operations. For instance, any part of a C program can interact with memory, the filesystem or the network. The program can even yield control to any location in program with the `goto` keyword, which has famously been criticized by Dijkstra (1968). This core of his argument is that `goto` breaks the structure of the code. The programmer then has to trace the execution of the program in their mind in order to understand it. The same reasoning extends to other effects too: the more effects a function is allowed to exhibit, the harder it becomes to reason about.

The “anything goes” approach to effects therefore puts a large burden of ensuring correct behaviour of effects on the programmer. If the language cannot give any guarantees about what (a part of) a program can do, the programmer has to check instead. Say, for instance, that a function somewhere in the code is modifying some global variable to some invalid value. This can then make entirely different parts of the program behave incorrectly. The programmer tasked with debugging this issue then has to examine the program as a whole to find where this modification takes place, if they even realize that it is set incorrectly. In languages where this is possible, effectful operations therefore limit our ability to split the code into chunks to be examined separately.

A solution is to treat certain effects in a more structured manner. For example, instead of allowing `goto`, a language might provide exceptions. In a language like Java, the exceptions are then also part of the type system, so that it is easier to track which functions are allowed to throw exceptions. However, this means that any effect must be backed by the language. That is, the language needs to have a dedicated feature for every effect that should be supported in this way and new effects cannot be created without adding a new feature to the language. This means that the support for various effects is always limited.

In contrast, languages adhering to the functional programming paradigm disallow effectful operations altogether.¹ Here, all functions are *pure*, meaning that they are functions in the mathematical sense: only a mapping from inputs to output. Such a function is *referentially transparent*, meaning that it always returns identical outputs for identical inputs. They also do not interact with their environment. By dictating that all functions are pure, a type

¹Usually there are some escape hatches to this rule, such as Haskell’s `trace` function, which is built-in and effectful, but only supposed to be used for debugging.

signature of a function becomes almost a full specification of what the function can do. In that sense, the return type is an all-encompassing description of what a function might do.

TODO: Example of manual state: checking balanced parentheses? Fibonacci? Reverse polish notation? Wadler & Peyton Jones use IO, that's nice I guess.

TODO: Pure means that the effect must be encoded in the return type of the function. We want types to be “all-encompassing”, “reliable”.

FEEDBACK:
vague

Yet this rule is quite limiting, since effects are often an important part of a computation. For instance, if we want to keep any mutable state `a` in a Haskell program, **we have to encode that state in the inputs and outputs of the program**. Manually threading the state through the program quickly becomes laborious in larger programs. The same goes for encodings of other effects. A more practical method of dealing with effectful operations in functional languages is through the use of *monads* (Peyton Jones and Wadler 1993; Wadler 1992).

Monads were introduced as a mathematical model for effectful computation by Moggi (1991). A function returning a monad is not fully executed. Instead, it is evaluated until the first effectful operation is encountered. This partially evaluated result is only further evaluated when it is passed to a *handler*. This handler decides what to do with this result and can resume the computation, which will again evaluate until the next effectful operation and the cycle repeats. If we then wish to have some state in our Haskell program, we have to wrap all our stateful functions in the `State` monad and pass it to the `runState` handler.

This split between the procedure and the handler provides some modularity. We can swap out the standard `runState` handler for some other handler. We might for instance write a handler that does not just yield the final value of the state, but a list containing the history of all values that the state has been set to. Or, we create a handler where every `get` operation increases the value of the state by one, such that every `get` yields a unique value. This is all possible without changing the code using the `State` monad.

The limitations of the monad approach become apparent when we look at procedures that use multiple effects. The problem is that the composition of two monads does not yield a monad. This is a limitation that can be worked around with *monad transformers*. A transformer can be applied to a monad to yield a new monad. In doing so, it adds more operations to the original monad. However, this yields complicated types which have to be taken into account while programming. Operations might need to be lifted from between monads. Additionally, monad transformers need **quadratic implementation stuffs**. Therefore, working with monad transformers is still quite laborious.

TODO: ADD EXAMPLES

As it turns out, however, not all effects are algebraic. *Higher-order effects* are effects with operations that take effectful computations as arguments that do not behave like continuations. The issue is that the handlers need be able to handle the effect and then let the rest of the computation evaluate, but that is not sufficient for higher-order effects.

In this thesis, we introduce a novel programming language called *Elaine*. The core idea of Elaine is to define a language which features elaborations and higher-order effects as a first-class construct. This brings the theory of hefty algebras into practice. With Elaine, we aim to demonstrate the usefulness of elaborations as a language feature. Throughout this thesis, we present example programs with higher-order effects to argue that elaborations are a natural and easy representation of higher-order effects.

Like handlers for algebraic effects, elaborations require the programmer to specify which elaboration should be applied. However, elaborations have several properties which make it likely that there is only one relevant possible elaboration. Hence, we argue that elaboration instead should often be implicit and inferred by the language. To this end, we introduce *implicit elaboration resolution*, a novel feature that infers an elaboration from the variables in scope.

Additionally, we give transformations from higher-order effects to algebraic effects. There are two reasons for defining such a transformation. The first is to show how elaborations can

FEEDBACK:
modularity is about syntax AND semantics. Reason about individual effects and compose them. Modular handlers with different implementations. And we can create new ‘syntax’

FEEDBACK:
 $n \times n$ problem?

be compiled in a larger compilation pipeline. The second is that these transformations show how elaborations could be added to existing systems for algebraic effects.

We present a specification for Elaine, including the syntax definition, typing judgments and semantics. Along with this specification, we provide a reference implementation written in Haskell in the artefact accompanying this thesis. This implementation includes a parser, type checker, interpreter, pretty printer, and the transformations mentioned above. Elaine opens up exploration for programming languages with higher-order effects. While not a viable general purpose language in its own right, it can serve as inspiration for future languages.

TODO: Overview of sections

1.1 Contributions

The main contribution of this thesis is the specification and implementation of Elaine. This consists of several parts.

- We define a syntax suitable for a language with both handlers and elaboration (Appendix B.1).
- We provide a set of examples for programming with higher-order effect operations.
- We present a type system for a language with higher-order effects and elaborations, based on Hindley-Milner type inference and inspired by the Koka type system. This type system introduces a novel representation of effect rows as multiset which, though semantically equivalent to earlier representations, allows for a simple definition of effect row unification.
- We propose that elaborations should be inferred in most cases and provide a type-directed procedure for this inference (Chapter 5).
- We define two transformations from programs with elaborations and higher-order effects to programs with only handlers and algebraic effects. The first transformation is convenient, but relies on impredicativity and therefore only works in languages that allow impredicativity, such as Elaine, Haskell and Koka. The second transformation is more involved, but does not rely on impredicativity and would therefore also be allowed in a language like Agda.

1.2 Artefact

TODO: Describe contents and structure of artefact

The artefact is available online at <https://github.com/tertsdiepraam/thesis/elaine>.

Chapter 2

Algebraic Effects

Elaine is based on the theory of hefty algebras, which is an extension of the theory of algebraic effects. Hence, the theory of algebraic effects also applies to Elaine. In this chapter, we give an introduction to algebraic effects. In the next chapter, we discuss its limitations regarding higher-order effects and describe how hefty algebras overcome those limitations.

2.1 Monads

TODO: Some more background and citations on monads

We will build up the notion of algebraic effects from monads. Monads are an abstraction over effectful computation commonly used in functional programming.

While many descriptions of monads using category theory and various analogies can be employed in explaining them, for our purposes, a monad is a type constructor `m` with two associated functions: **return** and `>>=`, with the latter pronounced “bind”. In Haskell, this concept is easily encoded in a type class, which is listed below.

```
1 class Monad m where
2   return :: a -> m a
3   (>>=) :: m a -> (a -> m b) -> m b
4
5   (>>) :: Monad m => m a -> m b -> m b
6   a >> b = a >>= \_ -> b
```

Haskell

This class tells us that we can construct a value of `m a` for any type `a` and for any monad `m` using **return**. Additionally, we can compose two monadic computations using `>>=`, which takes a monadic computation and a *continuation*, which is the function that should be called after the operation has been performed. The continuation is passed the return value of the operation as an argument.

To explain how effectful operations can be encoded with this, we can look at a simple example: the **Maybe** monad. Our goal with this monad is to create an “abort” effect, where the computation stops and returns immediately once **Nothing** is encountered.

```

1 data Maybe a
2   = Just a
3   | Nothing
4
5 class Monad Maybe where
6   return = Just
7
8   Just a  >>= k = k a
9   Nothing >>= k = Nothing

```

Haskell

With this definition, we can chain functions returning **Maybe**. For example, we can define a **head** function that returns the first element of a list if it is non-empty and **Nothing** otherwise. We can also define a division function which checks that the divisor is non-zero. These functions can then be composed using `>>=`.

```

1 head :: [a] -> Maybe a
2 head (x:xs) = Just x
3 head _ = Nothing
4
5 safeDiv :: Int -> Int -> Maybe Int
6 safeDiv _ 0 = Nothing
7 safeDiv x y = div x y
8
9 main = do
10   print $ head [] >>= safeDiv 10 -- -> Nothing
11   print $ head [0,1,2] >>= safeDiv 10 -- -> Nothing
12   print $ head [2,3,4] >>= safeDiv 10 -- -> Just 5

```

Haskell

A more involved example is the **State** monad. If we were to keep track of state manually a function that modifies state would need to take some state of type `s` as an argument and return a new value for the state. Therefore, if a function `foo` normally is a function with type `a -> b`, it would need to have the type `a -> s -> (s, b)`. Instead of modifying the state directly, it maps an old state to a new state. Then we need to ensure that we update the state with the modified value. For example, if the function is called multiple times, the code would look something like the code before.

```

1 -- Increment the state by `a` and return the old state
2 inc :: Int -> Int -> (Int, Int)
3 inc a s = (s + a, s)
4
5 multipleIncs :: Int -> (Int, Int)
6 multipleIncs s = let
7   (s', _) = inc 5 s
8   (s'', _) = inc 6 s'
9   in inc 7 s''

```

Haskell

The program becomes verbose and repetitive as a result. If we look at the signature `a -> s -> (s, b)`, we can see that there is a possibility for abstraction here: we can abstract over the pattern of `s -> (s, b)`. Our definition of the **State** type constructor then becomes:

```

1 | newtype State s a = State (s -> (s, a))
2 |
3 | -- so that the inc function becomes:
4 | inc :: Int -> State Int Int
5 | inc a = State $ \s -> (s + a, s)

```

Haskell

That might look like a step backwards at first, but we can now implement the monad functions for `State s` to allow us to compose functions returning the `State` type. Additionally, we define the `get` and `put` operations, which are the basic building blocks we can use to build more complex operations.

```

1 | instance Monad (State s) where
2 |   return x = State $ \s -> (s, x)
3 |
4 |   State fa >>= k = State $ \s ->
5 |     let (s', a) = fa s
6 |       State fb = k a
7 |     in fb s'
8 |
9 | get :: State s s
10 | get = State $ \s -> (s, s)
11 |
12 | put :: s -> State s ()
13 | put = State $ \s -> (s, ())

```

Haskell

For convenience, we also define `runState` to allow us to provide an initial state and evaluate the entire computation.

```

1 | runState :: s -> State s a -> (s, a)
2 | runState init (State s) = s init

```

Haskell

The `inc` operations can then be sequenced using the `>>=` operator. Because the return value of `inc` is irrelevant in the computation, we define a shorthand operator `>>`, which ignores the return value of the first operation.

```

1 | (>>) :: Monad m => m a -> m b -> m b
2 | ma >> mb = ma >>= \_ -> mb
3 |
4 | inc :: Int -> State Int Int
5 | inc x = get >>= \s -> put (s + x) >>= return s
6 |
7 | multipleIncs :: State Int Int
8 | multipleIncs = inc 5 >> inc 6 >> inc 7
9 |
10 | main = print $ runState 0 bar # prints 0 + 5 + 6 + 7 = 18

```

Haskell

This is the power of monads: they allow us to abstract the effectful operations away, while also signalling the effects that a function requires in the return type. In the final example, we do not have to think about how the `State` monad works anymore, but only use the `get` and `put` operations to build complex computations.

To make working with monads more convenient, Haskell also features `do`-notation, which is syntactic sugar for the `>>=` and `>>` operators. Using `do`-notation, the `multipleIncs` computation from the previous example can be written as

```
1 multipleIncs = do
2   inc 5
3   inc 6
4   inc 7
```

Haskell

If the results from the `inc` computations needs to be used, the `<-` operator, which is part of the `do`-notation, can be used to bind the result of a computation to a variable. For example, the sum of all the results from the `inc` calls can be returned.

```
1 multipleIncs = do
2   a <- inc 5
3   b <- inc 6
4   c <- inc 7
5   return (a + b + c)
6
7 -- which is equivalent to
8 multipleIncs =
9   inc 5 >>= \a ->
10    inc 6 >>= \b ->
11    inc 7 >>= \c ->
12    return (a + b + c)
```

Haskell

This is a convenient method for programming with effects in Haskell, while also staying true to its functional paradigm.

2.2 Effect Composition with the Free Monad

TODO: Cite Casper's blog or a more academic version of it from somewhere. <http://casperbp.net/posts/2023-07-algebraic-effects/index.html>

A limitation of the monads above is that they cannot be composed: it is not possible to make a computation with them that uses both `State` and `Maybe` at the same time. One solution to this is to use monad transformers, as explained in Section 7.1. Another solution is to use the *free monad*.

```
1 data Free f a
2   = Pure a
3   | Do (f (Free f a))
4
5 instance Functor f => Monad (Free f) where
6   return = Pure
7   Pure x >>= f = f x
8   Do g >>= f = Do $ fmap (>>= f) g
```

Haskell

As the listing above shows, the `Free` type constructor generates a monad for any given functor. So given some `State s` functor, then `Free (State s)` is a monad. Of course, this is only useful if the `State s` functor can generate a monad with the same functionality as the original state monad. To do so, we define a data type with the two constructors of `State`. This is a functor over the `k` parameter, which represents the *continuation* of the computation,

which is the rest of the computation to be evaluated after the effect operation. Note that we do not have to give definitions of **return** and **>>=** since those are provided by the free monad. We only have to derive the default **Functor** instance.

```
1 | data State s k = Put s k | Get (s -> k)
2 |   deriving Functor
```

Haskell

Similarly we can apply the free monad to **Maybe**. However, the **Just** constructor of the **Maybe** is already covered by the **Pure** constructor of the free monad, so **Maybe** can be simplified to a single constructor. We call this simplified type **Abort**. The **Abort** constructor does not use the continuation because it signals that evaluation should stop.

```
1 | data Abort k = Abort
2 |   deriving Functor
```

Haskell

In contrast with monads, can be meaningfully compose these functors. We define a type-level operator **+**, which can be thought of as **Either** for functors. We use this operator to build lists of functors. Just like lists have a **Cons** and **Nil**, these lists consist of **+** and **End**, where **End** is a functor without any constructor.

```
1 | infixr 6 +
2 | data (f + g) a = L (f a) | R (g a)
3 |   deriving Functor
4 |
5 | data End k
6 |   deriving Functor
```

Haskell

We can then make monads for any combination of the functors we have defined, such as **Free (State s + End)**, **Free (Abort + End)** or **Free (State s + Abort + End)**. In general, for any set of functors **f1**, ..., **fN**, we can construct a monad **Free (f1 + ... + fN + End)**.

However, we have no way to use any of the effect operations for this functor. For example, if we have **Free (State s + Abort + End)**, how would we use the **get** operation that we expect from the state monad? The solution is to give a definition for **get** for the free monad if and only if **State** is one of the composed functors. We do this with a typeclass relation **<**, which defines an injection from a functor **f** to any composed functor **g** that contains **f**. We can use this injection to define **get**, **put** and **abort**. These convenience functions are called *smart constructors*.

TODO: Parts of this might not be important


```

1  class f < g where
2      inj :: f k -> g k
3
4  instance f < f where inj = id
5  instance f < (f + g) where inj = L
6  instance f < h => f < (g + h) where inj = R . inj
7
8  get :: State s < f => Free f s
9  get = Do $ inj $ Get Pure
10
11 put  :: State s < f => s -> Free f ()
12 put s = Do $ inj $ Put $ Pure ()
13
14 abort :: Abort < f => Free f ()
15 abort = Do $ inj $ Abort

```

This makes it possible to construct a computation using all those operations. For example, a computation that checks the state, asserts that it is larger than 0 and then decrements the state by 1.

```

1  decrement :: Free (State Int + Abort + End) Int
2  decrement = get >>= \s ->
3      if s > 0
4      then put (s - 1) >>= pure (s - 1)
5      else abort

```

However, there is no way to evaluate this computation, because Haskell does not know what to do with these operations. To do that, there needs to be a function with the type

$$\text{Free } (f + f') \ a \rightarrow \text{Free } f' \ b$$

for every f and finally a $\text{Free } \text{End} \ a \rightarrow a$ to reduce the free monad to a final value. Following Plotkin and Pretnar (2009), these functions are called *handlers*. Using a fold over the free monad, the definition of the handlers can be reduced to two smaller functions:

- the return case, $a \rightarrow \text{Free } f' \ b$;
- and the case for handling f : $f \ (\text{Free } f' \ b) \rightarrow \text{Free } f' \ b$.

```

1 fold :: Functor f => (a -> b) -> (f b -> b) -> Free f a -> b
2 fold gen _ (Pure x) = gen x
3 fold gen alg (Op f)  = alg (fmap (fold gen alg) f)
4
5 handle :: (a -> Free f' b)
6         -> (f (Free f' b) -> Free f' b)
7         -> Free (f + f') a -> Free f' b
8 handle ret hndl = fold ret $
9   \case
10    L x -> hndl x
11    R x -> Do x
12
13 -- The Do case does not need to be handled since End cannot be
14 -- constructed
15 handleEnd :: Free End a -> a
16 handleEnd (Pure a) = a
17
18 handleAbort = handle (Pure . Just) (\Abort -> Nothing)

```

TODO: State handler

This finally allows us to use the `Abort` and `State` monads together, while providing a handler per functor. While the plumbing needed for a free monad is extensive, it is worth considering what it provides. First, we can combine multiple functors in our type signatures. Second, we can define operations that work for any effect composition that contains an effect. Third, we can provide modular handlers that handle a single effect from the composed functors. Any effect we define in this way is automatically compatible with all the other effects.

We have not only gained modularity for the effects themselves, but also for the handlers. There is nothing preventing different implementations of the handlers. It is, for example, possible to define a state handler in which `put` operations are ignored or in which state propagates “back in time”.

QUESTION:

Is it possible to write the state handler without `handle_` from Casper's blog post? I couldn't quite get the types to work, but that is basically what happens in Koka.

I saw this somewhere, but I forgot where, I want to say Wadler?

2.3 Algebraic Effects

The free monad encoding in the previous section is an implementation of algebraic effects in Haskell. The term “algebraic” comes from the fact that this method works for effects that can be described as algebraic theories (Plotkin and Power 2001). Later, Plotkin and Power (2003) showed that this is only possible for effects that satisfy the *algebraicity property*.

The algebraicity property states that the `>>=` operation distributes over the computation parameters of an operation. This can be derived from the definitions of `>>=` on `Free` and `fmap` for the effects. For example, we can apply the definitions to `Free (State s)`.

```

1 | -- we start with:
2 | put s >=> k
3 | -- expand smart constructor:
4 | Do (Put s (Pure ())) >=> k
5 | -- apply >=> of Free:
6 | Do $ fmap (>=> k) (Put s (Pure ()))
7 | -- apply fmap of Put:
8 | Do $ Put s (Pure () >=> k)
9 | -- simplify
10| Do $ Put s k

```

Haskell

Recalling the definition of `fmap`, the algebraicity property states that all constructor parameters of `k` need to be “continuation-like”. Their behaviour cannot change if the continuation is bound to it.

The state and abort effects satisfy this property, along with effects for non-determinism, cooperative concurrency. However, *higher-order effects* such as exception catching and the reader effect with a local operation are not algebraic. Those effects are discussed extensively in Chapter 3.

2.4 Building a Language with Algebraic Effects

Although the previous sections contain an encoding of algebraic effects, they are not yet very ergonomic to use. Since Haskell is designed around pure and monadic functions, it is not possible to hide all unnecessary details of the encoding. If we instead design a new language which integrates algebraic effects as a core feature in the language, we have much more freedom in designing a syntax and type system that work well for this purpose.

Elaine is a language with support for algebraic effects, but also for higher-order effects. Therefore, this section focuses on Koka (Leijen 2014, 2023), which only supports algebraic effects. Since Elaine is heavily inspired by Koka, the same concepts apply to Elaine.

At the core of such languages lies the following rule: all functions return the free monad. That means that a function `a -> b` is not allowed, but `a -> Free f b` is allowed. Because that is a bit verbose, we can introduce a shorthand: `a -> f b`. We will no longer refer to `f` as a functor, but as an *effect row* and its elements as *effects*. Because we change the name, we will use an `e` from now on instead of an `f` to denote an effect row. So, the function type `a -> e b` should be read as: this function takes an `a` and returns `b` with effect row `e`.

Instead of using type-level operators, we can introduce special syntax for effect rows, too. Following the lead of Koka (Leijen 2014), we will write effect rows as

$$\langle e_1, e_2, \dots, e_N \rangle.$$

In the type system, we are then allowed to use different orders of effects interchangeably. This is a clear ergonomic improvement over the free monad encoding, where we could only reason about inclusion of one effect at a time. With effect rows, the type reason about, for example, equality between $\langle a, b \rangle$ and $\langle b, a \rangle$. Additionally, it can distinguish $\langle \text{state}\langle a \rangle, \text{state}\langle a \rangle \rangle$ from $\langle \text{state}\langle a \rangle \rangle$, whereas with the functor interface we could only check `State s < f`, which would hold in both cases.

All the effects in this row are single effects, they are not composed. In Haskell, this is not the case, some functor `f` can represent a composed functor. Therefore we need notation to express that an effect row can be extended with another effect row. This is written as

$$\langle e_1, e_2, \dots, e_N | e_s \rangle,$$

where `es` is the tail of the effect row; a variable representing the effect row with which this effect row can be extended.

We can define the same effects as before, like state and abort, but in Koka, we do not define them as functors. Instead, we use the **effect** keyword. Each constructor of the functors then becomes a `ctl` operation.

```

1 effect abort
2   ctl abort(): a
3
4 effect state<a>
5   ctl get(): a
6   ctl put(x: a): ()

```

Koka

There are a few small differences. In Koka’s syntax, we cannot specify that `abort` does not resume, however, we can encode it in the type system instead. The return type `a` is a parameter of the function and, since we cannot construct an unknown type, it is functionally equivalent to a “never” or “empty” type, that is, a type that cannot be constructed. All operations in Koka are also functions, so `get` is a function here, too. However, we can make it a function that does not take any arguments.

With these types, the equivalent of `Free (State s + Abort + End) a` becomes `<state<s>,abort> a`. The equivalent of a handler would then be a function which takes `() -> <f|e> a` and returns `<|e> a`. In Koka, such a function can be defined with the **handler** construct, which requires an implementation for each operation of an effect and a special function for the return case. Note the similarity to the `handle` function we defined in Haskell before. In the case of abort effect, this handler is assigned to variable for later use. The state handler is wrapped in another function which takes an initial value for the state.

```

1 val hAbort = handler
2   return(x)   Just(x)
3   ctl abort() Nothing
4
5 fun hState(init, c)
6   val h = handler
7     return(x)   fn(s) x
8     ctl get()   fn(s) resume(s)(s)
9     ctl put(n)  fn(s) resume(())(n)
10
11   fn() h(c)(init)

```

Koka

This saves us from specifying some details, but the structure is largely the same as with the free monad encoding. The larger differences become apparent when we want to use the effects. A port of the decrement function is listed below.

```
1 fun decrement(): <state<int>,abort> int
2   val s = get()
3   if s == 0 then
4     abort()
5
6   put(s - 1)
7   s - 1
8
9 fun printMaybe(m: maybe<int>)
10  match m
11    Just(x) -> println(x)
12    Nothing -> println("nothing!")
13
14 fun main()
15   printMaybe(hAbort(hState(3, foo))) // prints "2"
16   printMaybe(hAbort(hState(0, foo))) // prints "nothing!"
```

Koka

The `>>=` operator is entirely implicit here. Therefore, it is similar to Haskell's `do`-notation. However, in `do`-notation, every effectful operation needs to be on a separate line. For example, if the state needs to be incremented by 1, this can be achieved in one line in Koka, but in Haskell using `do`-notation requires two lines.

```
1 put(get() + 1)
```

Koka

```
1 do
2   x <- get
3   put (x + 1)
```

Haskell

In Koka, effectful operations can be used anywhere as long as they are wrapped in a corresponding handler. In the end, the syntax here might even seem closer to imperative programming languages than functional programming languages. However, the type system is still very much like that of a functional language. For example, the type system is able to assert that a function is entirely pure. In the listing below, the `<>` in the type of the function asserts that it does not require effects, yet the `println` function requires an effect. Hence, Koka's type checker will yield a type error.

```
1 fun should_be_pure(x: int): <> int
2   println("This will give a type error!")
3   x + 10
```

Koka

As will become clear in Chapter 4, Elaine takes a lot of inspiration from Koka. Handlers and effects are defined in the same way, modulo some syntactical difference. What sets Elaine apart, is that it also supports higher-order effects, which will be explained in the next chapter.

Chapter 3

Higher-Order Effects

In the previous chapter, we explained the concept of algebraic effects; effects that satisfy the algebraicity property. We also mentioned that not all effects are algebraic. To be more specific, the effects that are not algebraic are higher-order effects: effects that take effectful computations as parameters. As a result, it is not possible to give modular implementations for these operations, like we can with algebraic effects. This chapter details the difficulties around higher-order effects and discusses hefty algebras, the theory that Elaine is based on.

3.1 Computation Parameters

Recall that an effect in the free monad encoding is a functor over some k with some constructors. The type k represents the continuation of the computation. Naturally, it is possible to write a constructor with multiple parameters of type k . For example, we could make a `Branch` functor which takes a boolean and two computations and selects the branch to take based on the boolean. It is essentially an **if-else** expression expressed as an effect.

```
1 data Branch k = Branch Bool k k
2
3 branch :: Branch < f => Bool -> Free f a -> Free f a -> Free f a
4 branch b ifTrue ifFalse :: Do $ inj $ Branch b ifTrue ifFalse
```

Haskell

The important observation with this effect is that both `ifTrue` and `ifFalse` behave like continuations. To examine why, consider the following computation.

```
1 branch b (pure 0) (pure 1) >>= \x -> pure (x + 1)
```

Haskell

Like previously established, the `>>=` operator distributes over the computation parameters. This yields the following expression.

```
1 branch b
2   (pure 0 >>= \x -> pure (x + 1))
3   (pure 1 >>= \x -> pure (x + 1))
4 -- which reduces to
5 branch b (pure 1) (pure 2)
```

Haskell

This computation has the same intended semantics as the original. The distribution of `>>=` therefore does not change the semantics and hence the effect is algebraic. Therefore, there would be no problem encoding this effect in Haskell using the encoding in the previous chapter and, by extension, in Koka.

This is what we mean by saying that the parameters are computation-like: the continuation can be appended to it without changing the semantics of the effect.

3.2 Breaking Algebraicity

For other effects, however, the intended semantics are not such that the computation parameters are continuation-like. These effects are called higher-order effects (citation needed).

One such effect is the **Reader** effect. Traditionally, the **Reader** monad has two operations: `local` and `ask`. The latter functions much like the `get` operation from the state effect and is therefore algebraic on its own. However, the `local` operation is more complex. It takes two parameters, a function `f` and a computation `c`. The intended semantics are then that whenever `ask` is used within `c`, the function `f` is applied to the returned value.

```
1 | data Reader a k = Ask (a -> k) | Local (a -> a) k k
2 |
3 | ask          = Do $ inj $ Get Pure
4 | local f c = Do $ inj $ Local f c (Pure ())
```

Haskell

To show why the `local` operation breaks algebraicity, consider the following computation.

```
1 | local (* 2) ask >>= \x -> ask >>= \y -> pure x + y
```

Haskell

Only the first `get` operation is inside the `local` operation and should therefore be doubled. If the **Reader** effect was algebraic, we should be able to distribute the `>>=` operator again without changing the semantics of the program. However, do that yield the following computation.

```
1 | local (* 2) (ask >>= \x -> ask >>= \y -> pure x + y)
```

Haskell

Now, both `get` operations are inside the `local` operation, so both values will be doubled. For example, if we had installed a handler that makes `ask` return 1, the first computation would return $2 + 1 = 3$ and the second $2 + 2 = 4$. Therefore, we have shown with a counterexample that the **Reader** effect cannot be algebraic.

A similar argument holds for the **Except** effect, which also has two operations: `catch` and `throw`. In the simplest form, `throw` resembles the **abort** effect, but it takes a value to abort the computation with. The `catch` operation should

```
1 | data Except a k = Throw a | Catch k k
2 |
3 | throw      = Do $ inj $ Throw
4 | catch a b = Do $ inj $ Catch a b
```

Haskell

Again, we take a simple example program to show that **Except** violates algebraicity.

```
1 | catch (pure False) (pure True) >>= throw -- -> throw False
2 | -- then distributing >>= yields
3 | catch (throw False) (throw True)          -- -> throw True
```

Haskell

Before distributing the `>>=` operator the computation should throw **False**, but after it should throw **True**. So, again, the semantics have changed by distributing the `>>=` and therefore **Except** is not algebraic.

3.3 The Modularity Problem

Taking a step back from effects, defining a function for exception catching is possible. Recall that the `throw` operation is algebraic, therefore, a handler for it can be defined. If we assume some handler for it called `handleThrow` which returns an **Either** where **Left** is the value from `throw` and **Right** is the value from a completed computation, we can define `catch` in terms of that function.

```

1 | catch c1 c2 =
2 |   case handleThrow c1 of
3 |     Left e -> c2
4 |     Right a -> pure a

```

Haskell

The distinction between effects which are and which are not algebraic has been described as the difference between *effect constructors* and *effect destructors* (Plotkin and Power 2003). The `local` and `catch` operations have to act on effectful computations and change the meaning of the effects in that computation. So, they have to deconstruct the effects in their computations using handlers. An imperfect heuristic for whether a function can be an algebraic effect is to check whether the implementation requires a handler. If it uses a handler, it probably cannot be an algebraic effect.

An algebraic effect can have a modular implementation: a computation can be reused in different contexts by using different handlers. For these higher-order effects such as `catch` and `local`, this is not possible. This is known as the *modularity problem* with higher-order effects (Wu, Schrijvers, and Hinze 2014). This is the motivation behind the research on higher-order effects, including this thesis. It is also the problem that the theory of hefty algebras aims to solve.

3.4 Hefty Algebras

Several solutions to the modularity problem have been proposed (van den Berg et al. 2021; Wu, Schrijvers, and Hinze 2014). Most recently, Bach Poulsen and van der Rest (2023) introduced hefty algebras. The idea behind hefty algebras is that there is an additional layer of modularity, specifically for higher-order effects.

For a full treatment of hefty algebras, we refer to Bach Poulsen and van der Rest (2023). In addition, the encoding of hefty algebras is explained in more detail by Bach Poulsen (2023).

At the core of hefty algebras are the hefty tree. A hefty tree is a generalization of the free monad to higher-order functors, which will write `HOFunctor`. In the listing below, we also repeat the definition of a functor from the previous chapter for comparison.

```

1 | -- A regular functor
2 | class Functor f where
3 |   fmap :: (a -> b) -> f a -> f b
4 |
5 | -- An higher-order functor
6 | class (forall f. Functor (h f)) => HOFunctor h where
7 |   hmap :: (f a -> g a) -> (h f a -> h g a)

```

Haskell

The definition of a hefty tree, with the free monad for reference, then becomes:


```

1  -- free monad
2  data Free f a
3    = Pure a
4    | Do (f (Free f a))
5
6  -- hefty tree
7  data Hefty h a
8    = Return a
9    | Do (h (Hefty h) (Hefty h a))

```

Haskell

A hefty tree and a free monad are very similar: we can define the $\gg=$, $<$ and $+$ operators from the previous chapter for hefty trees, so that the hefty tree can be used in the same way.¹ We refer to Bach Poulsen and van der Rest (2023) for the definition of these operators. Furthermore, any functor can be lifted to a higher-order functor with a `Lift` data type.

```

1  data Lift f (m :: * -> *) k = Lift (f k)
2    deriving Functor
3
4  instance Functor f => HOFunctor (Lift f) where
5    hmap _ (Lift x) = Lift x

```

Haskell

In algebraic effects, the evaluation of a computation can be thought of as a transformation of the free monad to the final result:

$$\text{Free } f \ a \xrightarrow{\text{handle}} b.$$

Using hefty algebras, the evaluation instead starts with a *hefty tree*, which is *elaborated* into the free monad. The full evaluation of a computation using hefty algebras then becomes:

$$\text{Hefty } h \ a \xrightarrow{\text{elaborate}} \text{Free } f \ a \xrightarrow{\text{handle}} b.$$

This elaboration is a transformation from a hefty tree into the free monad, defined as an algebra over hefty trees. The algebras are then used in `hfold`; a fold over hefty trees.

```

1  hfold :: HOFunctor h
2    => (forall a. a -> g a)
3    -> (forall a. h g (g a) -> g a)
4    -> Hefty h a
5    -> g a
6  hfold gen _ (Return x) = gen x
7  hfold gen alg (Do x)   =
8    ha alg (fmap (hfold gen alg) (hmap (hfold gen alg) x))
9
10 elab :: HOFunctor h
11    => (forall a. h (Free f) (Free f a) -> Free f a)
12    -> Hefty h a
13    -> Free f a
14 elab elabs = hfold Pure elabs

```

Haskell

For any algebraic – and thus lifted – effect, this elaboration is trivially defined by unwrapping the `Lift` constructor.

¹We are abusing Haskell’s syntax here. In the real Haskell encoding, these operators need to have different names from their free monad counterparts, for example :+ and :< .

```

1 | elabLift :: g < f => Lift g (Free f) (Free f a) -> Free f a
2 | elabLift (Lift x) = Op $ inj x

```

Haskell

Applying `elabLift` to `elab` then gives a function that elaborates `Hefty (Lift f) a` to `Free f a` for any functor `f`. The more interesting case is that of higher-order effects. For example, the `local` operation of the `Reader` effect can be mapped to a computation using the free monad as well, resembling the definition of `local` as a function.

```

1 | data Reader r k = Local r k k
2 |
3 | elabReader :: Ask r < f
4 |             => Reader r (Free f) (Free f a)
5 |             -> Free f a
6 | elabReader (Local f m k) = ask >>= \r -> handle (hAsk (f r)) m >>= k

```

Haskell

This definition of `elabReader` is modular, because it is a transformation of the computation. Even if the computation is fixed, the elaboration gives the `local` operation its meaning. Hence, hefty algebras solve the modularity problem.

These elaborations can be composed to construct elaborations for multiple effects as well. Bach Poulsen and van der Rest (2023) do this by introducing an operator which composes elaborations. However, that approach is limited because it requires all higher-order effects to be elaborated at a single location in the source code.

This limitation can be worked around by defining all operations such that they do not have type `Hefty h a`, but instead `(Hefty h a -> Free f a) -> Free f a`. The idea here is that every computation of that type is a partially elaborated computation, where the input is a function that specifies how to elaborate the remaining higher-order effects.

```

1 | type Comp h f a = (Hefty h a -> Free f a) -> Free f a
2 |
3 | local :: Reader r < h
4 |       => (r -> r)
5 |       -> Comp h f a
6 |       -> Comp h f a
7 |
8 | elab :: (Hefty h a -> Free f a)
9 |       -> Comp (h + h') f a
10 |      -> Comp h' f a
11 |
12 | elabEnd :: Hefty End a -> Free End a
13 |
14 | run :: Comp End f a -> Free f a

```

Haskell

TODO: Well this is turning out very ugly. It also requires `handle` to work correctly and change the `f` parameter. Alternatively, it would be to define elaborations more like handlers: `(Hefty (h + h') a -> Hefty h' a)`, but that kind of deviates more from the hefty algebras paper? I'm not sure what's best. Otherwise we could make a more informal argument?

Therefore, it is sound to define semantics where higher-order effects can be elaborated one by one, instead of all at once in a single `elab` call. This is one of the novel features that Elaine offers, as detailed in Section 4.7.

Chapter 4

A Tour of Elaine

The language designed for this thesis is called “Elaine”. The distinguishing feature of this language is its support for higher-order effects via elaborations. The basic feature of elaborations has been extended with two novel features: implicit elaboration resolution and compilation of elaborations, which are explained in Chapters 5 and 6, respectively.

This chapter introduces Elaine with motivating examples for the design choices. The full specification is given in Appendix B. Additional example programs are available in the artefact accompanying this thesis, the contents of which are detailed in Section 1.2 and in Appendix A.

4.1 Basics

As is tradition with introductions to programming languages, we have to start with a program that shows the string `"Hello, world!"`:

```
1 | # The value bound to main is the return value of the program
2 | let main = "Hello, world!";
```

Elaine

There are several aspects of Elaine that this example highlights. Comments start with `#` and continue until the end of the line. We bind variables with the `let` keyword. The `main` variable is required and is the value printed at the end of execution. In contrast with other languages, `main` is not a function in Elaine. Note that statements are required to end with a semicolon.

In addition to strings, Elaine features integers and booleans. The latter is expressed the `false` and `true` keywords. To operate on these types, we need functions to perform the operations. By default, there are no functions in scope, however, we can bring them in scope by importing the functions from the `std` module with the `use` keyword. For example, we can write a program that computes $5 \cdot 2 + 3$:

```
1 | use std;
2 | let main = add(mul(5, 2), 3);
```

Elaine

The `std` module contains functions for boolean and numeric arithmetic, comparison of values and more. The full list of functions is given in Appendix B.6. To show off some more functions, below is a program that returns a boolean indicating whether 2^5 is greater than $-(25 \cdot -30)$, which is true. Note that `-` is allowed as part of an integer literal, but not as an operator.

```
1 use std;
2 let main = gt(
3     pow(2, 5),
4     neg(mul(25, -30)),
5 );
```

Elaine

Bindings can be used to split up a computation, both at the top-level and within braces, which are used to group sequential expressions. A sequence of expressions evaluates to the last expression. Expressions are only allowed to contain variables that have been defined above the expression, so the order of bindings is significant. This rule also disallows recursion. Below is the same comparison written with some bindings.

```
1 use std;
2 let a = pow(2, 5);
3 let main = {
4     let b = mul(25, -30);
5     gt(a, neg(b))
6 };
```

Elaine

Functions are defined with **fn**, followed by a list of arguments and a function body. Unlike Haskell, functions are called with parentheses.

```
1 use std;
2 let double = fn(x) {
3     mul(2, x)
4 };
5 let square = fn(x) {
6     mul(x, x) # or pow(x, 2)
7 };
8 let main = double(square(8));
```

Elaine

Tuples are comma-separated lists of expressions surrounded with `()`. Tuples have a fixed length and can have elements of different types.

```
1 let main = (9, "hello");
```

Elaine

Additionally, Elaine features **if** expressions. The language does not support recursion or any other looping construct. Figure 4.1 contains a program that uses the basic features of Elaine and prints whether the square of 4 is even or odd.

4.2 Types

Elaine is strongly typed with Hindley-Milner style type inference. Let bindings, function arguments and function return types can be given explicit types. By convention, we will write variables and modules in lowercase and capitalize types.

The primitive types are `String`, `Bool` and `Int` for strings, booleans and integers respectively. We can specify the types for let bindings, function arguments and return types.

```

1 | # The standard library contains basic functions for manipulation
2 | # of integers, booleans and strings.
3 | use std;
4 |
5 | # Functions are created with `fn` and bound with `let`, just like
6 | # other values. The last expression in a function is returned.
7 | let square = fn(x: Int) Int {
8 |     mul(x, x)
9 | };
10 |
11 | let is_even = fn(x: Int) Bool {
12 |     eq(0, modulo(x, 2))
13 | };
14 |
15 | # Type annotations can be inferred:
16 | let square_is_even = fn(x) {
17 |     let result = is_even(square(x));
18 |     if result { "even" } else { "odd" }
19 | };
20 |
21 | let give_answer = fn(f, s, x) {
22 |     let prefix = concat(concat(s, " "), show_int(x));
23 |     let text = concat(prefix, " is ");
24 |     let answer = f(x);
25 |     concat(text, answer)
26 | };
27 |
28 | let main = give_answer(square_is_even, "The square of", 4);

```

Figure 4.1: A simple Elaine program. The result of this program is the string "The square of 4 is even".

```

1 | let x: Int = 5;      # ok!
2 | let x: String = 5;  # type error!
3 |
4 | let f = fn(x: Int) Int { mul(3, x) };
5 | let y = fn("Hello"); # type error!

```

We also could have written the type of the function as the type for the let binding. The type for a function is written like a function definition, without parameter names and body.

```

1 | let f: fn(Int) Int = fn(x) { mul(3, x) };

```

Type parameters start with a lowercase letter. They do not need to be declared explicitly, much like in Haskell.

```
1 | let f = fn(x: a) (a, Int) {  
2 |     (x, 5)  
3 | };  
4 | let y = f("hello");  
5 | let z = f(5);
```

Elaine

4.3 Algebraic Data Types

However, these primitive types do not allow us to write very complex programs. To do that, we need to be able to define our own data types. That is what the **type** construct is for. It is analogous to Koka's **type**, Haskell's **data** or Rust's **enum** construct.

A type consists of a list of constructors each with a list of parameters. These constructors can be used as functions. A type can have type parameters, which are declared with `[]` after the type name. It is not possible to put constraints on type parameters.

Data types can be deconstructed with the **match** construct. The **match** construct looks like Rust's **match** or Haskell's **case**, but is more limited. It can only be used for custom data types and only matches on the outer constructor. For example, it is not possible to match on `Just(5)`, but only on `Just(x)`.

```
1 | use std;  
2 |  
3 | type Maybe[a] {  
4 |     Just(a),  
5 |     Nothing(),  
6 | }  
7 |  
8 | let safe_div = fn(x, y) Maybe[Int] {  
9 |     if eq(y, 0) {  
10 |         Nothing  
11 |     } else {  
12 |         Just(div(x, y))  
13 |     }  
14 | };  
15 |  
16 | let main = match safe_div(5, 0) {  
17 |     Just(x) => show_int(x),  
18 |     Nothing => "Division by zero!",  
19 | };
```

Elaine

Since the `Maybe` type is very common, it is provided in the standard library under the `maybe` module.

Data types are allowed to be recursive. For example, we can define a `List` type as it is often defined in functional languages with a `Cons` and a `Nil` constructor.

```
1 | type List[a] {  
2 |     Cons(a, List[a]),  
3 |     Nil(),  
4 | }  
5 |  
6 | let list: List[Int] = Cons(1, Cons(2, Nil()));
```

Elaine

The `List` type is also provided in the standard library in the `list` module. If that module is in scope there is also some syntactic sugar for lists: we can write a list with brackets and comma-separated expressions like `[1, 2, 3]`.

4.4 Recursion, Loops and Lists

The `let` bindings in the previous sections are not allowed to be recursive. In general, `let` bindings can only reference values that have been defined before the binding itself. However, recursion or some other looping construct is necessary for many programs. Therefore, Elaine has a special syntax for recursive definitions: **`let rec`**.

`Let` bindings with **`rec`** definitions are desugared into the `Y` combinator. However, it is impossible to write the `Y` combinator manually, because it would have an infinite type. The type checker therefore has special case for recursive definitions.

An example of a recursive function is the `factorial` function listed below.

```
1 use std;
2
3 let rec factorial = fn(x: Int) {
4   if eq(x, 0) {
5     1
6   } else {
7     mul(x, factorial(sub(x,1)))
8   }
9 };
```

Elaine

A word of caution: Elaine has no guards against unbounded recursion of functions or even recursive expressions. For example, these statements are valid according to the Elaine type checker, but will cause infinite recursion when evaluated, until the interpreter runs out of memory:

```
1 # Warning: these declarations will not halt!
2 let rec f = fn(x) { f(x) };
3 let rec x = x;
```

Elaine

Using recursive definitions, we can build functions like `map`, `foldl` and `foldr` to operate on our previously defined `List` type. The implementation for `map` might look like the listing below.

```
1 let rec map = fn(f: fn(a) b, list: List[a]) List[b] {
2   match list {
3     Cons(a, as) => Cons(f(a), map(f, list)),
4     Nil() => Nil(),
5   }
6 };
7
8 let doubled = map(fn(x) { mul(2, x) }, [1, 2, 3]); # -> [2, 4, 6]
```

Elaine

Note that, in contrast with Haskell, Elaine evaluates `map` eagerly; there is no lazy evaluation.

The `list` module provides the most common operations on lists. Such `head`, `concat_list`, `range`, `map`, `foldl` and `foldr`. It also provides a `sum` function for lists of integers and a `join` function for lists of strings.

4.5 Algebraic Effects

The programs in the previous sections are all pure and contain no effects. While a monadic approach is possible, Elaine provides first class support for algebraic effects and effect handlers to make working with effects more ergonomic. The design of effects in Elaine is heavily inspired by Koka (Leijen 2014).

An effect is declared with the **effect** keyword. An effect needs a name and a set of operations. Operations are the functions that are associated with the effect. They can have an arbitrary number of arguments and a return type. Only the signature of operations can be given in an effect declaration, the implementation must be provided via handlers (see Section 4.5.1).

Figure 4.2 lists examples of effect declarations for the **Abort**, **Ask**, **State** and **Write** effects. We will refer to those declarations throughout this section. For the listings in this section, one can assume that these declarations are used. The **Abort** effect is meant to exit the computation. **Ask** provides some integer value to the computation, much like a global constant. **State** corresponds to the **State** monad in Haskell. Finally, **Write** allows us to write some string value somewhere. We will be using this to provide a substitute for writing to standard output.

<pre> 1 effect Abort { 2 abort() (), 3 }</pre>	<pre> 1 effect Ask { 2 ask() Int, 3 }</pre>
<pre> 1 effect State { 2 get() Int, 3 put(Int) (), 4 }</pre>	<pre> 1 effect Write { 2 write(String) (), 3 }</pre>

Figure 4.2: Examples of algebraic effect declarations for some simple effects.

4.5.1 Effect Handlers

To define the implementation of an effect, we have to define a handler it. Handlers are first-class values in Elaine and can be created with the **handler** keyword. They can then be applied to an expression with the **handle** keyword. When **handle** expressions are nested with handlers for the same effect, the innermost **handle** applies.

For example, if we want to use an effect to provide an implicit value, we can make an effect **Ask** and a corresponding handler, which **resumes** execution with some values. The **resume** function represents the continuation of the program after the operation. The simplest handler for **Ask** we can write is one which yields some constant value.

```

1 | let hAsk = handler { ask() { resume(10) } };
2 |
3 | let main = handle[hAsk] add(ask(), ask()); # evaluates to 20
```

Elaine

Of course, it would be cumbersome to write a separate handler for every value we would like to provide. Since handlers are first-class values, we can return the handler from a function to simplify the code. This pattern is quite common to create dynamic handlers with small variations.

```

1 | let hAsk = fn(v: Int) {
```

Elaine

```

2   handler { ask() { resume(v) } }
3 };
4
5 let main = {
6   let a = handle[hAsk(6)] add(ask(), ask());
7   let b = handle[hAsk(10)] add(ask(), ask());
8   add(a, b)
9 };

```

The true power of algebraic effects, however, lies in the fact that we can also write a handler with entirely different behaviour, without modifying the computation. For example, we can create a stateful handler which increments the value returned by `ask` on every call to provide unique identifiers. The program below will return 3, because the first `ask` call returns 1 and the second returns 2. This is accomplished in a very similar manner to the State monad.

```

1 let hAsk = handler {
2   return(x) { fn(s: Int) { x } }
3   ask() {
4     fn(s: Int) {
5       let f = resume(s);
6       f(add(s, 1))
7     }
8   }
9 };
10
11 let c = handle[hAsk] add(ask(), ask());
12 let main = c(1);

```

Elaine

Calling the `resume` function is not required. All effect operations are executed by the `handle` expression, hence, if we return from the operation, we return from the `handle` expression.

The `Abort` effect is an example which does not call the continuation. It defines a single operation `abort`, which stops the evaluation of the computation. The canonical handler for `Abort`, which returns the `Maybe` monad. If the computation returns, it should wrap the returned value in `Just`. Otherwise, if the computation aborts, it should return `Nothing()`. In Elaine, if a sub-computation of a handler returns, the optional `return` arm of the handler will be applied. In the code below, this wraps the returned value in a `Just`. All arms of a handler must have the same return type.

```

1 effect Abort {
2   abort() ()
3 }
4
5 let hAbort = handler {
6   return(x) { Just(x) }
7   abort() { Nothing() }
8 };
9
10 let main = handle[hAbort] {
11   abort();
12   5

```

Elaine

```
13 |};
```

Alternatively, we can define a handler that defines a default value for the computation in case it aborts. This is more convenient than the first handler if the `abort` case should always become

```
1  let hAbort = fn(default) {
2      handler {
3          return(x) { x }
4          abort() { default }
5      }
6  };
7
8  let safe_div = fn(x, y) <Abort> Int {
9      if eq(y, 0) {
10         abort()
11     } else {
12         div(x, y)
13     }
14 };
15
16 let main = add(
17     handle[hAbort(0)] safe_div(3, 0),
18     handle[hAbort(0)] safe_div(10, 2),
19 );
```

Elaine

Just like we can ignore the continuation, we can also call it multiple times, which is useful for non-determinism and logic programming. In the listing below, the `Twice` effect is introduced, which calls its continuation twice. Combining that with the `State` effect as previously defined, the `put` operation is called twice, incrementing the initial state 3 by two, yielding a final result of 5. Admittedly, this example is a bit contrived. A more useful application of this technique can be found in Appendix A.1, which contains the full code for a very naive SAT solver in Elaine, using multiple continuations.

```

1  effect Twice {
2      twice() ()
3  }
4
5  let hTwice = handler {
6      twice() {
7          resume(());
8          resume(())
9      }
10 }
11
12 let main = {
13     let a = handle[hState] handle[hTwice] {
14         twice();
15         put(add(get(), 1));
16         get()
17     };
18     a(3)
19 };

```

Elaine

4.5.2 Effect Rows

All types in Elaine have an effect row. Up to this point, we have omitted the effect rows. This works because effect rows are inferred by the type checker. Effect rows represent the set of effects that need to be handled to obtain the value in a computation. For simple values, that effect row is empty, denoted $\langle \rangle$. For example, an integer has type $\langle \rangle$ `Int`. Without row elision, the `square` function in the previous section could therefore have been written as

```

1  let square = fn(x:  $\langle \rangle$  Int)  $\langle \rangle$  Int {
2      mul(x, x)
3  };

```

Elaine

Simple effect rows consist of a list of effect names separated by commas. The return type of a function that returns an integer and uses the `Ask` and `State` effects has type $\langle \text{Ask}, \text{State} \rangle$ `Int` or, equivalently $\langle \text{State}, \text{Ask} \rangle$ `Int`. The order of effects in effect rows is irrelevant. However, the multiplicity is important, that is, the effect rows $\langle \text{State}, \text{State} \rangle$ and $\langle \text{State} \rangle$ are not equivalent. To capture the equivalence between effect rows, we therefore model them as multisets.

Additionally, we can extend effect rows with other effect rows. In the syntax of the language, this is denoted with the `|` at the end of the effect row: $\langle A, B | e \rangle$ means that the effect row contains `A`, `B` and some (possibly empty) set of remaining effects. We called a row without extension *closed* and a row with extension *open*.

Like types, effect rows are unified in the type checker. For unification any closed row is first opened by introducing a new expansion variable. Then unification solves for the equation

$$\langle A_1, \dots, A_N | e \rangle = \langle B_1, \dots, B_M | f \rangle,$$

for e and f . To do so, a fresh variable g is introduced which represents the intersection of e and f . The unified row then becomes $\langle A_1, \dots, A_N, B_1, \dots, B_M | g \rangle$. Table 4.1 provides some more examples of effect row unification. The full procedure for unification is detailed in Appendix B.2.

$\langle A \rangle$	\cup	$\langle \rangle$	\rightarrow	$\langle A \rangle$
$\langle A \rangle$	\cup	$\langle A \rangle$	\rightarrow	$\langle A \rangle$
$\langle A \rangle$	\cup	$\langle B \rangle$	\rightarrow	$\langle A, B \rangle$
$\langle A, B \rangle$	\cup	$\langle B, A \rangle$	\rightarrow	$\langle A, B \rangle$
$\langle A, A \rangle$	\cup	$\langle A \rangle$	\rightarrow	$\langle A, A \rangle$
$\langle A, B e \rangle$	\cup	$\langle C \rangle$	\rightarrow	$\langle A, B, C e' \rangle$
$\langle A e \rangle$	\cup	$\langle B e \rangle$	\rightarrow	$\# \text{ error!}$
$\langle A e1 \rangle$	\cup	$\langle B e2 \rangle$	\rightarrow	$\langle A, B e3 \rangle$

Table 4.1: Examples of effect row unification.

We can use extensions to ensure equivalence between effect rows without specifying the full rows. For example, the following function uses the `Abort` effect if the called function returns false, while retaining the effects of the wrapped function.

```
1 | let abort_on_false = fn(f: fn()  $\langle$ |e $\rangle$  Bool)  $\langle$ Abort|e $\rangle$  () {
2 |     if f() { () } else { abort() }
3 | }
```

Elaine

When an effect is handled, it is removed from the effect row. The `main` binding is required to have an empty effect row, which means that all effects in the program need to be handled. Therefore, to use the `abort_on_false` function defined above, it needs to be called from within a handler.

```
1 | let main:  $\langle$  $\rangle$  Maybe[()] = handle[hAbort] {
2 |     abort_on_false(fn() { false })
3 | };
```

Elaine

4.6 Effect-Agnostic Functions

Recall the definition of `map` in Section 4.4, which was written without any effects in its signature. Adding the effects yields the following definition.

```
1 | pub let rec map = fn(f: fn(a)  $\langle$ |e $\rangle$  b, l: List[a])  $\langle$ |e $\rangle$  List[b] {
2 |     match l {
3 |         Nil() => Nil(),
4 |         Cons(x, xs) => Cons(f(x), map(f, xs)),
5 |     }
6 | };
```

Elaine

Note that the parameter `f` and `map` use the same effect row variable `e`. This means that `map` has the same effect row as `f` for any effect row that `f` might have, including the empty effect row. This makes `map` quite powerful, because it can be applied in many situations.

```
1 | let pure_doubled = map(fn(x) { mul(2, x) }, [1,2,3]);
2 | let ask_added = handle[hAsk(5)] map(fn(x) { add(ask() x) }, [1,2,3]);
```

Elaine

If we were to write the same expressions in Haskell instead, we would need two different implementations of `map`: one for applying pure functions (`map`) and another for applying monadic functions (`mapM`). Our definition of `map` is therefore more general than Haskell's `map`.

function. The same reasoning can be applied to other functions like `foldl` and `foldr` or indeed any higher-order function.

4.7 Higher-Order Effects

Higher-order effects in Elaine are supported via elaborations, as proposed by Bach Poulsen and van der Rest (2023) and explained in Section 3.4. In this framework, higher-order effects are elaborated into a computation using only algebraic effects. They are not handled directly. This means that we cannot write handlers for them as we did for algebraic effects in the previous section.

To distinguish higher-order effects and operations from algebraic effects and operations, we write them with a `!` suffix. For example, a higher-order `Exception!` effect is written `Exception!`, and its `catch` operations is written `catch!`.

Higher-order effects are treated exactly like algebraic effects in the effect rows. The order of effects still does not matter, and we can create effect rows with arbitrary combinations of algebraic and higher-order effects.

The elaborated operations differ from other functions and algebraic operations because they have call-by-name semantics; the arguments are not evaluated before they are passed to the elaboration. Hence, the arguments can be computations, even effectful computations.

Just like we have the `handler` and `handle` keywords to create and apply handlers for algebraic effects, we can create and apply elaborations with the `elaboration` and `elab` keywords. Unlike handlers, elaborations do not get access to the `resume` function, because they always resume exactly once.

An illustrative example of this feature is the `Reader` effect with a `local` operation, shown in `??`. This effect enhances the previously introduced `Ask` effect with a `local` operation that modifies the value returned by `ask`. To motivate the implementation in `??`, let us first imagine how to emulate the behaviour of `local`. Our goal is to make the following snippet return the value 15.

```
1 | let main = handle[hAsk(5)] {
2 |   let x = ask();
3 |   let y = local(double, fn() { ask() });
4 |   add(x, y)
5 | };
```

Elaine

This means that the `local` operation would need to handle the `ask` effect with the modified value. This is easily achieved, since the innermost handler always applies. If the function to modify the value is called `f`, then the value we should provide to the handler is `f(ask())`.

```
1 | let local = fn(f: fn(Int) Int, g: fn() <Ask|e> a) <|e> a {
2 |   handle[hAsk(f(ask()))] { g() }
3 | }
```

Elaine

This works but is not implemented as an effect. For example, we cannot modularly provide another implementation of `local`. To turn this implementation into an effect, we start with the effect declaration.

```
1 | effect Reader! {
2 |   local!(fn(Int) Int, a) a
3 | }
```

Elaine

It might be surprising that the signature of `local` does not match the signature of the function above. That is because of the call-by-name nature of higher-order operations: instead of a function returning `a`, we simply have a computation that will evaluate to `a`. The effect row is irrelevant and therefore implicit. Now we can provide an elaboration, which is not a function, but better described as a syntactic substitution.

```
1 let eLocal = elaboration Reader! -> <Ask> {
2   local!(f, c) {
3     handle[hAsk(f(ask()))] c
4   }
5 }
```

Elaine

Note how similar the elaboration for `local!` is to the `local` function above. In the first line, we specify explicitly what effect the elaboration elaborates (`Reader!`) and which effects should be present in the context where this elaboration is used (`<Ask>`). This can be an effect row of multiple effects if necessary. In this case we only require the `Ask` effect. This means that we can use this elaboration in any expression that is wrapped by at least a handler for `Ask`.

```
1 let main = handle[hAsk(5)] elab[eLocal] {
2   let x = ask();
3   let y = local!(double, ask());
4   add(x, y)
5 }
```

Elaine

That is the full implementation for the higher-order `Reader!` effect in Elaine. Appendix A.2 contains a listing of all these pieces put together in a single example.

Another example is the `Exception!` effect. This effect should allow us to use the `catch!` operation to recover from a `throw`. The latter is an algebraic, so we can start there.

```
1 type Result[a, b] {
2   Ok(a),
3   Err(b),
4 }
5
6 effect Throw {
7   throw(String) a
8 }
9
10 let hThrow = handler {
11   return(x) { Ok(x) }
12   throw(s) { Err(s) }
13 };
```

Elaine

We assume here that we want to throw some error message, but we could put a different type in there as well. The `throw` operation has a return type `a`, which is impossible to construct in general, so it cannot return. The higher-order `Exception!` effect should then look like this:

```

1 | effect Exception! {
2 |     throw!(String) a
3 |     catch!(a, a) a
4 | }

```

Elaine

In contrast with the `Reader!` effect above, we alias the operation of the underlying algebraic effect here. This makes no functional difference, except that it allows us to write functions with explicit effect rows with `Exception!` and without `Throw`. We might even choose to elaborate to a different effect than `Throw`. The downside is that it requires us to provide the elaboration for the `throw!` operation.

```

1 | let eExcept = elaboration Exception! -> <Throw> {
2 |     throw!(s) { throw(s) }
3 |     catch!(a, b) {
4 |         match handle[hThrow] a {
5 |             Ok(x) => x,
6 |             Err(s) => b,
7 |         }
8 |     }
9 | };

```

Elaine

We can then use the effect like we used the `Reader!` effect: with an **elab** for `Exception!` and a **handle** for `Throw`. In the listing below, we ensure that we do not decrement a value of `0` to ensure it will not become negative.

```

1 | let main = handle[hThrow] elab[eExcept] {
2 |     let x = 0;
3 |     catch!({
4 |         if eq(x, 0) {
5 |             throw!("Whoa, x can't be zero!")
6 |         } else {
7 |             sub(x, 1)
8 |         }
9 |     }, 0)
10 | };

```

Elaine

Since the elaborations can be swapped out, we can also design elaborations with different behaviour. Assume, for instance, that there is a `Log` effect. Then we can create an alternative elaboration that logs the errors it catches, which might be useful for debugging.


```
1 | let eExceptLog = elaboration Exception! -> <Throw,Log> {  
2 |   throw!(s) { throw(s) }  
3 |   catch!(a, b) {  
4 |     match handle[hThrow] a {  
5 |       Ok(x) => x,  
6 |       Err(s) => {  
7 |         log(s);  
8 |         b  
9 |       }  
10 |    }  
11 |  }  
12 | };
```

Elaine

We could also disable exception catching entirely if we so desire. This might be helpful if we are debugging a piece of a program that is wrapped in a `catch!` to ensure it never fully crashes, but we want to see errors while we are debugging. Of course, this changes the functionality of the program significantly. We should therefore be careful not to change computations that rely on a specific implementation of the `Exception!`.

```
1 | let eExceptIgnoreCatch = elaboration Exception! -> <Throw> {  
2 |   throw!(s) { throw(s) }  
3 |   catch!(a, b) { a }  
4 | }
```

Elaine

What these examples illustrate is that elaborations provide a great deal of flexibility, with which we can define and alter the functionality of the `Exception!` effect. We can change it temporarily for debugging purposes or apply another elaboration to a part of a computation. We can also define more `Exception`-like effects and use multiple at the same time.

Chapter 5

Implicit Elaboration Resolution

With Elaine, we aim to explore further ergonomic improvements for programming with effects. We note that elaborations are often not parametrized and that there is often only one in scope at a time. Hence, when we encounter an **elab**, there is only one possible elaboration that could be applied. Therefore, we propose that the language should be able to infer the elaborations. Take the example in the listing below, where we let Elaine infer the elaboration for us.

```
1 let eLocal = elaboration Reader! -> <Ask> {
2   local!(f, c) {
3     handle[hAsk(f(ask()))] c
4   }
5 };
6
7 let main = handle[hAsk(2)] elab {
8   local!(double, add(ask(), ask()));
9 };
```

Elaine

A use case of this feature is when an effect and elaboration are defined in the same module. When this module is imported, the effect and elaboration are both brought into scope and **elab** will apply the standard elaboration automatically.

```
1 mod local {
2   pub effect Ask { ... }
3   pub let hAsk = handler { ... }
4   pub effect Reader! { ... }
5   pub let eLocal = elaboration Reader! -> <Ask> { ... }
6 }
7
8 use local;
9
10 # We do not have to specify the elaboration, since it is
11 # imported along with the effect.
12 let main = handle[hAsk] elab { local!(double, ask!()) };
13 #
```

Elaine

However, while useful, this feature only saves a few characters in the examples above. It becomes more important when multiple higher-order effects are involved: and **elab** without argument will elaborate all higher-order effects in the sub-computation. For instance, if elaborations for both **Exception** and **Reader** are in scope, the following program works.

```

1 | let main = handle[hAsk(2)] handle[hThrow] elab {
2 |     local!(double, {
3 |         if gt(ask(), 3) {
4 |             throw()
5 |         } else {
6 |             add(ask(), 4)
7 |         }
8 |     })
9 | }

```

Elaine

This relies on the fact that the order in which elaboration are applied does not affect the semantics of the program as explained in Section 3.4. To make the inference predictable, we require that an implicit elaboration must elaborate all higher-order effects in the sub-computation.

A problem with this feature arises when multiple elaborations for a single effect are in scope; which one should then be used? To keep the result of the inference predictable and deterministic, the type checker should yield a type error in this case. Hence, if type checking succeeds, then the inference procedure has found exactly one elaboration to apply for each higher-order effect. If not, the elaboration cannot be inferred and must be written explicitly.

```

1 | let eLocal1 = elaboration Local! -> <Ask> { ... };
2 | let eLocal2 = elaboration Local! -> <Ask> { ... };
3 |
4 | let main = elab { local!(double, ask!()) }; # Type error here!

```

Elaine

The elaboration resolution consists of two parts: inference and transformation. The inference is done by the type checker and is hence type-directed, which records the inferred elaboration. After type checking the program is then transformed such that all implicit elaborations have been replaced by explicit elaborations.

To infer the elaborations, the type checker first analyses the sub-expression. This will yield some computation type with an effect row containing both higher-order and algebraic effects: $\langle H1!, \dots, HN!, A1, \dots, AM \rangle$. It then checks the type environment to look for elaborations $e1, \dots, eN$ which elaborate $H1!, \dots, HN!$, respectively. Only elaborations that are directly in scope are considered, so if an elaboration resides in another module, it needs be imported first. For each higher-order effect, there must be exactly one elaboration.

The **elab** is finally transformed into one explicit **elab** per higher-order effect. Recall that the order of elaborations does not matter for the semantics of the program, meaning that we apply them in arbitrary order.

A nice property of this transformation is that it results in very readable code. Because the elaboration is in scope, there is an identifier for it in scope as well. The transformation then simply inserts this identifier. The **elab** in the first example of this chapter will, for instance, be transformed to **elab**[eVal]. A code editor could then display this transformed **elab** as an inlay hint.

Chapter 6

Elaboration Compilation

Since Elaine has a novel semantics for elaborations, it is worth examining its relation to well-studied constructs from programming language theory. Therefore, we introduce a transformation from programs with higher-order effects to a program with only algebraic effects, translating higher-order effects into algebraic effects, while preserving their semantics.

The goal of this transformation is twofold. First, it further connects hefty algebras and Elaine to existing literature. For example, by compiling to a representation with only algebraic effects, we can then further compile the program using existing techniques, such as the compilation procedures defined for Koka (Leijen 2017). In this thesis and the accompanying implementation, we provide the first step of this compilation. Second, the transformation allows us to encode elaborations in existing libraries and languages for algebraic effects.

6.1 Non-locality of Elaborations

TODO: Actually it's not just non-locality but also undecidable

Examining the semantics of elaborations, we observe that elaborations perform a syntactic substitution. For instance, the program on the left transforms into the program on the right by replacing `plus_two!`, with the expression `{ x + 2 }`.

<pre>1 use std; 2 3 effect PlusTwo! { 4 plus_two!(Int) 5 } 6 7 let ePlusTwo = { 8 elaboration PlusTwo! -> <> { 9 plus_two!(x) { add(x, 2) } 10 } 11 }; 12 13 let main = elab plus_two!(5);</pre>	<div>Elaine</div> <div>Elaine</div>
<pre>1 let main = { add(5, 2) };</pre>	<div>Elaine</div>

Additionally, the location of the `elab` does not matter as long as the operations are evaluated within it. For instance, these expressions are equivalent:

<pre>1 let main = elab[e] { 2 a!(); 3 a!() 4 };</pre>	<div>Elaine</div>
<pre>1 let main = { 2 elab[e] a!(); 3 elab[e] a!() 4 };</pre>	<div>Elaine</div>

In some cases, it is therefore possible to statically determine the elaboration that should be applied. In that situation, we can remove the elaboration from the program by performing the syntactic substitution.

However, we cannot apply that technique in general. One example where it does not work is when the elaboration is given by a complex expression, such as an **if**-expression:

```
1 | elab[if cond { elab1 } else { elab2 }] c
```

Elaine

Moreover, a single operation might need to be elaborated by different **elab** constructs, depending on run-time computations. In the listing below, there are two elaborations **eOne** and **eTwo** of an operation **a!()**. The **a!()** operation in **f** is elaborated where **f** is called. If the condition **k** evaluates to **true**, **f** is assigned to **g**, which is elaborated by **eOne**. However, if **k** evaluates to **false**, **f** is called in the inner **elab** and hence **a!()** is elaborated by **eTwo**.

```
1 | elab[eOne] {
2 |   let g = elab[eTwo] {
3 |     let f = fn() { a!() };
4 |     if k {
5 |       f
6 |     } else {
7 |       f();
8 |       fn() { () }
9 |     }
10 |   }
11 |   g()
12 | }
```

Elaine

Therefore, the analysis of determining the elaboration that should be applied to an operation is non-local. The static substitution could be used as an optimization or simplification step, but it cannot guarantee that the transformed program will not contain higher-order effects.

6.2 Operations as Functions

As explained in Appendix B.5, higher-order operations are evaluated differently from functions. The main difference is that the arguments are thunked and passed by name, instead of by value.

This behaviour can be emulated for functions if anonymous functions are passed as arguments instead of expressions. That is, for any operation call **op!(e1, ..., eN)**, we wrap the arguments into functions to get **op(fn() { e1 }, ..., fn() { eN })**. In the body of **op**, we then replace each occurrence of an argument **x** with **x()** such that the thunked value is obtained. The **op** operation can then be evaluated like a function instead, but it still has the intended semantics.

6.3 Compiling Elaborations to Dictionary Passing

TODO: This is currently just an example. The actual transformation needs to be clearly defined too.

Instead, the elaborations can be transforms with a technique similar to dictionary-passing style: the implicit context of elaborations is explicitly passed to functions that require a higher-order effect.

Listing 6.1: Untransformed program. This example should use a higher-order effect.

```

1  use std;
2
3  effect A! {
4      arithmetic!(Int, Int) Int
5  }
6
7  let eAdd = elaboration A! -> <> {
8      arithmetic!(a, b) { add(a, b) }
9  };
10
11 let eMul = elaboration A! -> <> {
12     arithmetic!(a, b) { mul(a, b) }
13 };
14
15 let foo = fn(k) {
16     elab[eAdd] {
17         let g = elab[eMul] {
18             let f = fn() { a!(5, 2 + 3) };
19             if k {
20                 f
21             } else {
22                 let x = f();
23                 fn() { x }
24             }
25         };
26         g()
27     }
28 };

```

Any function with higher-order effects then takes the elaboration to apply as an argument and the operation is wrapped in an **elab**. The elaboration is then taken from **elabA** at the call-site.

An example of what this transformation is given in listings 6.1 and 6.2, where listing 6.1 shows an untransformed program with higher-order effects and listing 6.2 shows the result of the transformation.

6.4 Compiling Elaborations into Handlers

TODO: Talk about impredicativity and how that makes it so that it does not work.

While the transformation in the previous section is correct, the transformed program is quite verbose, because the elaboration types need to be passed to every function with higher-order effects. It would be more convenient if this was passed implicitly.

As it turns out, we have a mechanism for passing implicit arguments: algebraic effects! Conceptually, both **elab** and **handle** are similar: they define a scope in which a given elaboration or handler is used. This scope is the same for both.

To use this observation, we start by defining a handler that returns an elaboration for higher-order effect **A!**, much like the **Ask** effect from Chapter 4.

Listing 6.2: Transformed program after compiling elaborations to dictionary passing.

```
1  use std;
2
3  effect A! {
4      arithmetic!(Int, Int) Int
5  }
6
7  let eAdd = elaboration A! -> <> {
8      arithmetic!(a, b) { add(a, b) }
9  };
10
11 let eMul = elaboration A! -> <> {
12     arithmetic!(a, b) { mul(a, b) }
13 };
14
15 # Create a new type with one constructor to represent the
16 # elaboration. The fields of the constructor are the
17 # operations. We assume for convenience that all the
18 # generated identifiers do not conflict with existing
19 # identifiers.
20 type ElabA {
21     ElabA(fn(fn() Int, fn() Int) Int)
22 }
23
24 # Convenience function to access the operation a from A!
25 let elab_A_a = fn(e: ElabA) {
26     let ElabA(v) = e;
27     v
28 };
29
30 let eAdd = ElabA ( fn(a, b) { add(a(), b()) } );
31 let eMul = ElabA ( fn(a, b) { mul(a(), b()) } );
32
33 # The transformed program
34 let foo = fn(k) {
35     let elab_A = eAdd;
36     let g = {
37         let elab_A = eMul;
38         let f = fn(elab_A) {
39             elab_A_a(elab_a)(fn() { 5 }, fn() { 2 + 3 })
40         };
41         if k {
42             f
43         } else {
44             let x = f(elab_A);
45             fn(elab_A) { x }
46         }
47     };
48     g(elab_A_a)
49 };
```

Combining the ideas above, we obtain a surprisingly simple transformation. Each elaboration is transformed into a handler, which resumes with a function containing the original expression, where argument occurrences force the thunked values. Since elaborations are now handlers, we need to change the **elab** constructs to **handle** constructs accordingly. Finally, the arguments to operation calls are thunked and the function that is resumed is called, that is, there is an additional $()$ at the end of the operation call.

$$\begin{array}{lcl}
 \mathbf{elab}[\$e_1\$] \{ \$e_2\$ \} & \Longrightarrow & \mathbf{handle}[\$e_1\$] \{ \$e_2\$ \} \\
 \\
 \begin{array}{l}
 \mathbf{elaboration} \{ \\
 \quad \$op_1!(x_{\{1,1\}}, \dots, x_{\{k_1,1\}})\$ \{ \$e_1\$ \} \\
 \quad \dots \\
 \quad \$op_n!(x_{\{1,n\}}, \dots, x_{\{k_n,n\}})\$ \{ \$e_n\$ \} \\
 \}
 \end{array} & \Longrightarrow & \begin{array}{l}
 \mathbf{handler} \{ \\
 \quad \$op_1(x_{\{1,1\}}, \dots, x_{\{k_1,1\}}) \\
 \quad \quad \text{resume}(\mathbf{fn}() \{ \$e_1\$[\$x_{\{1,1\}}] \}) \\
 \} \\
 \quad \dots \\
 \quad \$op_n(x_{\{1,n\}}, \dots, x_{\{k_n,n\}}) \\
 \quad \quad \text{resume}(\mathbf{fn}() \{ \$e_1\$[\$x_{\{1,n\}}] \}) \\
 \} \\
 \}
 \end{array} \\
 \\
 \$op_j\$!(\$e_1, \dots, e_k\$) & \Longrightarrow & \$op_j\$(\mathbf{fn}() \backslash \{ \$e_1\$ \backslash \}, \dots, \mathbf{fn}())
 \end{array}$$

The simplicity of the transformation makes it alluring and begs the question: are dedicated language features for higher-order effects necessary or is a simpler approach possible?

TODO: Answer: yes if you care about impredicativity, no otherwise.

Chapter 7

Related Work

Ever, since the definition of algebraic effects, it has been clear that not all effects could be encoded as algebraic effects (Plotkin and Power 2001). In the meantime, various theories have been proposed to overcome this limitation and solve the modularity problem for higher-order effects. This chapter discusses some of these theories and the libraries and languages that have been built with them.

Additionally, this chapter details some other languages that support algebraic effects in different ways from Elaine.

FEEDBACK: Can be expanded. Provide context for this thesis. What have others done, what's missing, and what does this thesis add?

7.1 Monad Transformers

Monad transformers provide a way to compose monads (Moggi 1989). This makes them an alternative to the free monad. While monad transformers predate algebraic effects, they do support higher-order effects.

The goal of monad composition is to make the operations of all composed monads available to the computation. Given two monads A and B , a naive composition would result in the type $A (B\ a)$. However, this type represents a computation using A that returns a computation $B\ a$, meaning that it is not possible to use operations of both monads.

A monad transformer is a type constructor that takes some monad and returns a new monad. Usually, the transformation it performs is to add operations to the input monad. Composing A and B then requires some transformer AT to be defined, such $AT\ B$ is a monad that provides the operations of both A and B . An arbitrary number of monad transformers can be composed this way. The representation of a monad then becomes much like that of a list of monad transformers. The `Identity` monad marks the end of the list, and it is defined as

```
1 | newtype Identity a = Identity a
```

Haskell

A popular implementation of monad transformers is Haskell's `mtl` library. In the rest of this section, we adopt the terminology from that library.

A neat property of monad transformers is that a monad can be easily obtained by applying the transformer to the identity monad. Haskell's `mtl` library, for instance, defines a monad transformer `StateT` and then defines `State` as `StateT Identity`.

The operations of the state effect are then not implemented on `StateT` directly, but on are part of a type class `MonadState`. The `StateT` is then an instance of `MonadState` class. Every other transformer is an instance of `MonadState` if its input monad is an instance of `MonadState`. For example, for the `WriterT` instance, there is the following instance declaration.

```

1 instance MonadState s (StateT s m) where
2   -- definitions omitted
3
4 instance MonadState s m => MonadState s (WriteT m) where
5   -- definitions omitted

```

Haskell

A computation can then be generic over the monad transformers, requiring only that `StateT` is present somewhere in the stack of monad transformers.

```

1 usesState :: MonadState Int m => Int -> m Int
2 usesState a = get >>= \x -> put (x + a)

```

Haskell

This is analogous to the `State s < f` constraint from the free monad encoding. However, there is a cost to this approach. For every effect, a new type class needs to be introduced and there need to be instance definitions on all existing monad transformers. The number of instance declarations therefore scales quadratically with the number of effects.

Another downside to monad transformers is that the order in which the monads need to be evaluated is entirely fixed. In the free monad encoding and languages with algebraic effects, the effects in the effect row can be reordered. **The order of the monad transformers also determines the order in which they must be handled: the outermost monad transformer must be handled first.**

In practice, this model has turned out to work quite well, especially in combination with `do`-notation, which allowed for easier sequential execution of effectful computations.

TODO: Contextual vs parametric effect rows (see effects as capabilities paper). The paper fails to really connect the two: contextual is just parametric with implicit variables. However, it might be more convenient. The main difference is in the interpretation of purity (real vs contextual). In general, I'd like to have a full section on effect row semantics. In the capabilities paper effect rows are sets, which makes it possible to do stuff like (Leijen 2005).

As the theoretical research around effects has progressed, new libraries and languages have emerged using the state-of-the-art effect theories. These frameworks can be divided into two categories: effects encoded in existing type systems and effects as first-class features.

These implementations provide ways to define, use and handle effectful operations. Additionally, many implementations provide type level information about effects via *effect rows*. These are extensible lists of effects that are equivalent up to reordering. The rows might contain variables, which allows for *effect row polymorphism*.

7.2 Effects as Free Monads

There are many libraries that implement the free monad in various forms in Haskell, including `fused-effects`¹, `polysemy`², `freer-simple`³ and `eff`⁴. Each of these libraries give the encoding of effects a slightly different spin in an effort to find the most ergonomic and performant representation. They are all not just based on the free monad, but on freer monads (Kiselyov and Ishii 2016) and fused effects (Wu and Schrijvers 2015) for better performance.

Effect rows are often constructed using the *Data Types à la Carte* technique (Swierstra 2008), which requires a fairly robust type system. Hence, many languages cannot encode effects within the language itself. In some languages, it is possible to work around the limitations with metaprogramming, such as the Rust library `effin-mad`⁵, though the result does

¹<https://github.com/fused-effects/fused-effects>

²<https://github.com/polysemy-research/polysemy>

³<https://github.com/lexi-lambda/freer-simple>

⁴<https://github.com/hasura/eff>

⁵<https://github.com/rosefromthedead/effin-mad>

not integrate well with the rest of language and its use in production is strongly discouraged by the author.

The programming language Idris (Brady 2013) also has an implementation of algebraic effects in its standard library. It is an interesting case study since Idris is a dependently typed language. Due to its dependent typing, it can distinguish multiple occurrences of a single effect in the same effect row by assigning them different *labels*. This is similar to what *named handlers* (Xie et al. 2022) aims to accomplish.

Some of these libraries support *scoped effects* (Wu, Schrijvers, and Hinze 2014), which is a limited but practical frameworks for higher-order effects. It can express the `local` and `catch` operations, but some higher-order effects are not supported.

Add λ -abstraction as example.

7.3 Prior Art for First-Class Effects

The motivation of adding support for effects to a programming language is twofold. First, it enables effects to be implemented into languages with type systems in which effects cannot be encoded as a free monad or a similar model. Second, built-in effects allow for more ergonomic and performant implementations. Naturally, the ergonomics of any given implementation are subjective, but we undeniably have more control over the syntax by adding effects to the language.

Notable examples of languages with first-class support for algebraic effects are Eff (Bauer and Pretnar 2015), Koka (Leijen 2014), OCaml (citation needed), and Frank (Lindley, McBride, and McLaughlin 2017).

In all of these languages, effect row variables can be used to abstract over effects. For example, the signature of the `map` function in Koka is given below and is similar to the signature of `map` in Elaine.

```
1 | fun map ( xs : list<a>, f : a -> e b ) : e list<b>
2 |   ...
```

Koka

Other languages choose a more implicit syntax for effect polymorphism. Frank (Lindley, McBride, and McLaughlin 2017) opts to have the empty effect row represent the *ambient effects*. Any effect row then becomes not the exact set of effects that need to be handled, but the smallest set. The equivalent signature of `map` is then written as

```
1 | map : {X -> []Y} -> List X -> []List Y
```

Frank

In contrast with Elaine, languages such as Koka and Frank do not have dedicated types for handlers and **handle** constructs. Instead, they represent handlers as functions that take computations as arguments. In Elaine, there are dedicated types and constructs for effect handlers so that they are symmetric with elaborations. That is, the counterpart of **elab** is **handle** and the counterpart of **elaboration** is **handler**.

Several extensions to algebraic effects have been explored in the languages mentioned above. Koka supports scoped effects and named handlers (Xie et al. 2022), which provides a mechanism to distinguish between multiple occurrences of an effect in an effect row. In Koka, scoped effects are effects for which a scope variable is created for every handler, which can be used to construct types that can only be used within the handler and cannot escape it. Note that this notion of scoped effects is different from the notion of scoped effects by Piróg et al. (2018), Wu, Schrijvers, and Hinze (2014), and Yang et al. (2022).

Chapter 8

Conclusion

TODO:

Bibliography

- Bach Poulsen, Casper (2023). *Algebras of Higher-Order Effects in Haskell*. URL: <http://casperbp.net/posts/2023-08-algebras-of-higher-order-effects/> (visited on 09/09/2023).
- Bach Poulsen, Casper and Cas van der Rest (Jan. 9, 2023). “Hefty Algebras: Modular Elaboration of Higher-Order Algebraic Effects”. In: *Proceedings of the ACM on Programming Languages* 7 (POPL), pp. 1801–1831. ISSN: 2475-1421. DOI: 10.1145/3571255. URL: <https://dl.acm.org/doi/10.1145/3571255> (visited on 01/26/2023).
- Bauer, Andrej and Matija Pretnar (Jan. 2015). “Programming with algebraic effects and handlers”. In: *Journal of Logical and Algebraic Methods in Programming* 84.1, pp. 108–123. ISSN: 23522208. DOI: 10.1016/j.jlamp.2014.02.001. URL: <https://linkinghub.elsevier.com/retrieve/pii/S2352220814000194> (visited on 04/03/2023).
- Van den Berg, Birthe et al. (2021). “Latent Effects for Reusable Language Components”. In: *Programming Languages and Systems*. Ed. by Hakjoo Oh. Vol. 13008. Series Title: Lecture Notes in Computer Science. Cham: Springer International Publishing, pp. 182–201. ISBN: 978-3-030-89050-6 978-3-030-89051-3. DOI: 10.1007/978-3-030-89051-3_11. URL: https://link.springer.com/10.1007/978-3-030-89051-3_11 (visited on 01/12/2023).
- Brady, Edwin (Sept. 25, 2013). “Programming and reasoning with algebraic effects and dependent types”. In: *Proceedings of the 18th ACM SIGPLAN international conference on Functional programming*. ICFP’13: ACM SIGPLAN International Conference on Functional Programming. Boston Massachusetts USA: ACM, pp. 133–144. ISBN: 978-1-4503-2326-0. DOI: 10.1145/2500365.2500581. URL: <https://dl.acm.org/doi/10.1145/2500365.2500581> (visited on 06/13/2023).
- Dijkstra, Edsger W. (Mar. 1968). “Letters to the editor: go to statement considered harmful”. In: *Communications of the ACM* 11.3, pp. 147–148. ISSN: 0001-0782, 1557-7317. DOI: 10.1145/362929.362947. URL: <https://dl.acm.org/doi/10.1145/362929.362947> (visited on 05/31/2023).
- Kiselyov, Oleg and Hiromi Ishii (Jan. 28, 2016). “Freer monads, more extensible effects”. In: *ACM SIGPLAN Notices* 50.12, pp. 94–105. ISSN: 0362-1340, 1558-1160. DOI: 10.1145/2887747.2804319. URL: <https://dl.acm.org/doi/10.1145/2887747.2804319> (visited on 06/03/2023).
- Leijen, Daan (July 23, 2005). “Extensible records with scoped labels”. In: — (June 5, 2014). “Koka: Programming with Row Polymorphic Effect Types”. In: *Electronic Proceedings in Theoretical Computer Science* 153, pp. 100–126. ISSN: 2075-2180. DOI: 10.4204/EPTCS.153.8. URL: <http://arxiv.org/abs/1406.2061v1> (visited on 06/16/2023).
- (Jan. 2017). “Type directed compilation of row-typed algebraic effects”. In: *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages*. POPL ’17: The 44th Annual ACM SIGPLAN Symposium on Principles of Programming Languages.

- Paris France: ACM, pp. 486–499. ISBN: 978-1-4503-4660-3. DOI: 10.1145/3009837.3009872. URL: <https://dl.acm.org/doi/10.1145/3009837.3009872> (visited on 04/08/2023).
- Leijen, Daan (Mar. 15, 2023). *The Koka Programming Language*. URL: <https://koka-lang.github.io/koka/doc/book.html> (visited on 06/13/2023).
- Lindley, Sam, Conor McBride, and Craig McLaughlin (Oct. 3, 2017). *Do be do be do*. arXiv: 1611.09259[cs]. URL: <http://arxiv.org/abs/1611.09259> (visited on 04/08/2023).
- Moggi, Eugenio (1989). *An Abstract View of Programming Languages*.
- (July 1991). “Notions of computation and monads”. In: *Information and Computation* 93.1, pp. 55–92. ISSN: 08905401. DOI: 10.1016/0890-5401(91)90052-4. URL: <https://linkinghub.elsevier.com/retrieve/pii/0890540191900524> (visited on 01/12/2023).
- Peyton Jones, Simon L. and Philip Wadler (1993). “Imperative functional programming”. In: *Proceedings of the 20th ACM SIGPLAN-SIGACT symposium on Principles of programming languages - POPL ’93*. the 20th ACM SIGPLAN-SIGACT symposium. Charleston, South Carolina, United States: ACM Press, pp. 71–84. ISBN: 978-0-89791-560-1. DOI: 10.1145/158511.158524. URL: <http://portal.acm.org/citation.cfm?doid=158511.158524> (visited on 06/03/2023).
- Piróg, Maciej et al. (July 9, 2018). “Syntax and Semantics for Operations with Scopes”. In: *Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science. LICS ’18: 33rd Annual ACM/IEEE Symposium on Logic in Computer Science*. Oxford United Kingdom: ACM, pp. 809–818. ISBN: 978-1-4503-5583-4. DOI: 10.1145/3209108.3209166. URL: <https://dl.acm.org/doi/10.1145/3209108.3209166> (visited on 09/11/2023).
- Plotkin, Gordon and John Power (2001). “Adequacy for Algebraic Effects”. In: *Foundations of Software Science and Computation Structures*. Ed. by Furio Honsell and Marino Miculan. Red. by Gerhard Goos, Juris Hartmanis, and Jan van Leeuwen. Vol. 2030. Series Title: Lecture Notes in Computer Science. Berlin, Heidelberg: Springer Berlin Heidelberg, pp. 1–24. ISBN: 978-3-540-41864-1 978-3-540-45315-4. DOI: 10.1007/3-540-45315-6_1. URL: http://link.springer.com/10.1007/3-540-45315-6_1 (visited on 04/08/2023).
- (2003). “Algebraic Operations and Generic Effects”. In: *Applied Categorical Structures* 11.1, pp. 69–94. ISSN: 09272852. DOI: 10.1023/A:1023064908962. URL: <http://link.springer.com/10.1023/A:1023064908962> (visited on 06/03/2023).
- Plotkin, Gordon and Matija Pretnar (2009). “Handlers of Algebraic Effects”. In: *Programming Languages and Systems*. Ed. by Giuseppe Castagna. Vol. 5502. Series Title: Lecture Notes in Computer Science. Berlin, Heidelberg: Springer Berlin Heidelberg, pp. 80–94. ISBN: 978-3-642-00589-3 978-3-642-00590-9. DOI: 10.1007/978-3-642-00590-9_7. URL: http://link.springer.com/10.1007/978-3-642-00590-9_7 (visited on 04/08/2023).
- Swierstra, Wouter (July 2008). “Data types à la carte”. In: *Journal of Functional Programming* 18.4. ISSN: 0956-7968, 1469-7653. DOI: 10.1017/S0956796808006758. URL: http://www.journals.cambridge.org/abstract_S0956796808006758 (visited on 06/11/2023).
- Wadler, Philip (1992). “The essence of functional programming”. In: *Proceedings of the 19th ACM SIGPLAN-SIGACT symposium on Principles of programming languages - POPL ’92*. the 19th ACM SIGPLAN-SIGACT symposium. Albuquerque, New Mexico, United States: ACM Press, pp. 1–14. ISBN: 978-0-89791-453-6. DOI: 10.1145/143165.143169. URL: <http://portal.acm.org/citation.cfm?doid=143165.143169> (visited on 08/24/2023).
- Wu, Nicolas and Tom Schrijvers (2015). “Fusion for Free: Efficient Algebraic Effect Handlers”. In: *Mathematics of Program Construction*. Ed. by Ralf Hinze and Janis Voigtländer. Vol. 9129. Series Title: Lecture Notes in Computer Science. Cham: Springer International Publishing, pp. 302–322. ISBN: 978-3-319-19796-8 978-3-319-19797-5. DOI: 10.1007/978-3-319-19797-5_15. URL: https://link.springer.com/10.1007/978-3-319-19797-5_15 (visited on 09/11/2023).

Wu, Nicolas, Tom Schrijvers, and Ralf Hinze (June 10, 2014). *Effect Handlers in Scope*.

Xie, Ningning et al. (Oct. 31, 2022). “First-class names for effect handlers”. In: *Proceedings of the ACM on Programming Languages* 6 (OOPSLA2), pp. 30–59. ISSN: 2475-1421. DOI: 10.1145/3563289. URL: <https://dl.acm.org/doi/10.1145/3563289> (visited on 06/13/2023).

Yang, Zhixuan et al. (2022). “Structured Handling of Scoped Effects: Extended Version”. In: Publisher: arXiv Version Number: 1. DOI: 10.48550/ARXIV.2201.10287. URL: <https://arxiv.org/abs/2201.10287> (visited on 06/13/2023).

Appendix A

Elaine Example Programs

This chapter contains longer Elaine samples with some additional explanation.

A.1 A naive SAT solver

This program is a naive brute-forcing SAT solver. We first define a `Yield` effect, so we can yield multiple values from the computation. We will use this to find all possible combinations of boolean inputs that satisfy the formula. The `Logic` effect has two operations. The `branch` operation will call the continuation twice; once with **false** and once **true**. With `fail`, we can indicate that a branch has failed. To find all solutions, we just `branch` on all inputs and `yield` when a correct solution has been found and `fail` when the formula is not satisfied. In the listing below, we check for solutions of the equation $\neg a \wedge b$.

```
1 use std;
2
3 effect Yield {
4   yield(String) ()
5 }
6
7 effect Logic {
8   branch() Bool
9   fail() a
10 }
11
12 let hYield = handler {
13   return(x) { "" }
14   yield(m) {
15     concat(concat(m, "\n"), resume(()))
16   }
17 };
18
19 let hLogic = handler {
20   return(x) { () }
21   branch() {
22     resume(true);
23     resume(false)
24   }
25   fail() { () }
26 };
```

Elaine

```
27
28 let show_bools = fn(a, b, c) {
29   let a = concat(show_bool(a), ", ");
30   let b = concat(show_bool(b), ", ");
31   concat(concat(a, b), show_bool(c))
32 };
33
34 let f = fn(a, b, c) { and(not(a), b) };
35
36 let assert = fn(f, a, b, c) <Logic,Yield> () {
37   if f(a, b, c) {
38     yield(show_bools(a, b, c))
39   } else {
40     fail()
41   }
42 };
43
44 let main = handle[hYield] handle[hLogic] {
45   assert(f, branch(), branch(), branch());
46 };
```

A.2 The Reader Effect

```
1 use std;
2
3 effect Ask {
4   ask() Int
5 }
6
7 effect Reader! {
8   local!(fn(Int) Int, a) a
9 }
10
11 let hAsk = fn(v: Int) {
12   handler {
13     return(x) { x }
14     ask() { resume(v) }
15   }
16 };
17
18 let eLocal = elaboration Reader! -> <Ask> {
19   local!(f, c) {
20     handle[hAsk(f(ask()))] c
21   }
22 };
23
24 let double = fn(x) { mul(2, x) };
25
26 let main = handle[hAsk(2)] elab[eLocal] {
27   local!(double, add(ask(), ask()));
28 };
```

Elaine

Appendix B

Elaine Specification

TODO: Custom type declarations are in the language but not explained in this chapter yet.

This chapter contains the detailed specification for Elaine: the syntax, semantics, the type inference rules and finally some specifics on the type checker that deviate from standard Hindley-Milner type checking.

B.1 Syntax definition

The Elaine syntax was designed to be relatively easy to parse. The grammar is white-space insensitive and most constructs are unambiguously identified with keywords at the start.

Based on the previous chapters, the `elab` without an elaboration might be surprising. The use of that syntax is explained in Chapter 5.

The full syntax definition is given in Figure B.1. For convenience, we define and use several extensions to BNF:

- tokens are written in **monospace font**, this includes the tokens `[]`, `<>`, `|` and `!`, which might be confused with the syntax of BNF,
- `[p]` indicates that the sort `p` is optional,
- `p...p` indicates that the sort `p` can be repeated zero or more times, and
- `p, ..., p` indicates that the sort `p` can be repeated zero or more times, separated by commas.

B.2 Effect row semantics

Before explaining the typing judgments of Elaine, let us examine effect rows. The effect row of a computation type determines the context in which the computation can be evaluated. For example, a computation with effect row `<A,B,C>` is valid in a function with effect row `<A,B,C>`. Additionally, the effect rows `<A,B>` and `<B,A>` should be considered to be equivalent.

One possible treatment is then to model effect rows as sets. However, as noted by Leijen (2014), this leads to some problems. Consider the following (abridged) program.

```
1 let v: fn(f: fn() <abort|e> a) e a {  
2   handle[hAbort] f()  
3 };  
4  
5 let main = handle[hAbort] v(fn() { abort() });
```

Elaine

$$\begin{aligned}
&\text{program } p ::= d \dots d \\
&\text{declaration } d ::= [\text{pub}] \text{ mod } x \{d \dots d\} \\
&\quad | [\text{pub}] \text{ use } x; \\
&\quad | [\text{pub}] \text{ let } p = e; \\
&\quad | [\text{pub}] \text{ effect } \phi \{s, \dots, s\} \\
&\quad | [\text{pub}] \text{ type } x \{s, \dots, s\} \\
&\text{block } b ::= \{ es \} \\
&\text{expression list } es ::= e; es \\
&\quad | \text{let } p = e; es \\
&\quad | e \\
&\text{expression } e ::= x \\
&\quad | () \mid \text{true} \mid \text{false} \mid \text{number} \mid \text{string} \\
&\quad | \text{fn}(p, \dots, p) [T] b \\
&\quad | \text{if } e \text{ b else } b \\
&\quad | e(e, \dots, e) \mid \phi(e, \dots, e) \\
&\quad | \text{handler } \{\text{return}(x) \text{ b}, o, \dots, o\} \\
&\quad | \text{handle}[e] e \\
&\quad | \text{elaboration } x! \rightarrow \Delta \{o, \dots, o\} \\
&\quad | \text{elab}[e] e \mid \text{elab } e \\
&\quad | es \\
&\text{annotatable variable } p ::= x : T \mid x \\
&\text{signature } s ::= x(T, \dots, T) T \\
&\text{effect clause } o ::= x(x, \dots, x) b \\
&\text{type } T ::= \Delta \tau \mid \tau \\
&\text{value type } \tau ::= x \\
&\quad | () \mid \text{Bool} \mid \text{Int} \mid \text{String} \\
&\quad | \text{fn}(T, \dots, T) T \\
&\quad | \text{handler } x \tau \tau \\
&\quad | \text{elaboration } x! \Delta \\
&\text{effect row } \Delta ::= \langle \phi, \dots, \phi[|x] \rangle \\
&\text{effect } \phi ::= x \mid x!
\end{aligned}$$

Figure B.1: Syntax definition of Elaine

The function v “removes” an **abort** effect from the effect row. By treating the effect row as a set, there would be no **abort** effect in return type of v . However, in **main**, there is another handler for **abort** and hence **abort** should be in the effect row.

The treatment of effect rows then simplifies if duplicated effects are allowed (Leijen 2014). Hence, we use multisets to model effect rows, meaning that the row $\langle A, B, B, C \rangle$ is represented by the multiset $\{A, B, B, C\}$. This yields a semantics where the multiplicity of effects is significant, but the order is not.

Since the effect row of a computation must match the effect row of the context in which it is used, the effect row of the computation is an overapproximation of the effects that are necessary. Therefore, we should allow effect row polymorphism, so that the same expression can be used within multiple contexts.

Effect row polymorphism is enabled via the *row tail*, which is denoted with the $|$ symbol followed by an identifier.

The $|$ symbol signifies extension of the effect row with another (possibly arbitrary) effect row. We determine compatibility between effect rows by unifying them. That is

We define the operation set as follows:

$$\begin{aligned} \text{set}(\varepsilon) &= \text{set}(\langle \rangle) = \emptyset \\ \text{set}(\langle A_1, \dots, A_n \rangle) &= \{A_1, \dots, A_n\} \\ \text{set}(\langle A_1, \dots, A_n | R \rangle) &= \text{set}(\langle A_1, \dots, A_n \rangle) + \text{set}(R). \end{aligned}$$

Note that the extension uses the sum, not the union of the two sets. This means that $\text{set}(\langle A | \langle A \rangle \rangle)$ should yield $\{A, A\}$ instead of $\{A\}$.

Then we get the following equality relation between effect rows A and B :

$$A \cong B \iff \text{set}(A) = \text{set}(B).$$

In typing judgments, the effect row is an overapproximation of the effects that actually used by the expression. We freely use set operations in the typing judgments, implicitly calling the set function on the operands where required. An omitted effect row is treated as an empty effect row $\langle \rangle$.

Any effect prefixed with a $!$ is a higher-order effect, which must be elaborated instead of handled. Due to this distinction, we define the operations $H(R)$ and $A(R)$ representing the higher-order and first-order subsets of the effect rows, respectively. The same operators are applied as predicates on individual effects, so the operations on rows are defined as:

$$H(\Delta) = \{\phi \in \Delta \mid H(\phi)\} \quad \text{and} \quad A(\Delta) = \{\phi \in \Delta \mid A(\phi)\}.$$

TODO: Talk about (Leijen 2005, 2014).

During type checking effect rows are represented as a pair consisting of a multiset of effects and an optional extension variable. In this section we will use a more explicit notation than the syntax of Elaine by using the multiset representation directly. Hence, a row $\langle A_1, \dots, A_n | e_A \rangle$ is represented as the multiset $\{A_1, \dots, A_n\} + e_A$.

Like with regular Hindley-Milner type inference, two rows can be unified if we can find a substitution of effect row variables that make the rows equal. For effect rows, this yields 3 distinct cases.

If both rows are closed (i.e. have no extension variable) there are no variables to be substituted, and we just employ multiset equality. That is, to unify rows A and B we check that $A = B$. If that is true, we do not need to unify further and unification has succeeded. Otherwise, we cannot make any substitutions to make them equal and unification has failed.

If one of the rows is open, then the set of effects in that row need to be a subset of the effects in the other row. To unify the rows

$$A + e_A \quad \text{and} \quad B$$

we assert that $A \subseteq B$. If that is true, we can substitute e_n for the effects in $B - A$.

Finally, there is the case where both rows are open:

$$A + e_A \quad \text{and} \quad B + e_B.$$

In this case, unification is always possible, because both rows can be extended with the effects of the other. We create a fresh effect row variable e_C with the following substitutions:

$$\begin{aligned} e_A &\rightarrow (B - A) + e_C \\ e_B &\rightarrow (A - B) + e_C. \end{aligned}$$

In other words, A is extended with the effects that are in B but not in A and similarly, B is extended with the effects in A but not in A .

B.3 Typing judgments

The context $\Gamma = (\Gamma_M, \Gamma_V, \Gamma_E, \Gamma_\Phi)$ consists of the following parts:

$\Gamma_M : x \rightarrow (\Gamma_V, \Gamma_E, \Gamma_\Phi)$	module to context
$\Gamma_V : x \rightarrow \sigma$	variable to type scheme
$\Gamma_E : x \rightarrow (\Delta, \{f_1, \dots, f_n\})$	higher-order effect to elaboration type
$\Gamma_\Phi : x \rightarrow \{s_1, \dots, s_n\}$	effect to operation signatures

INFO: A Γ_T for data types might be added.

Whenever one of these is extended, the others are implicitly passed on too, but when declared separately, they not implicitly passed. For example, Γ'' is empty except for the single $x : T$, whereas Γ' implicitly contains Γ_M, Γ_E & Γ_Φ .

$$\Gamma'_V = \Gamma_V, x : T \quad \Gamma''_V = x : T$$

If the following invariants are violated there should be a type error:

- The operations of all effects in scope must be disjoint.
- Module names are unique in every scope.
- Effect names are unique in every scope.

B.3.1 Type inference

We have the usual generalize and instantiate rules. But, the “generalize” rule requires an empty effect row.

QUESTION: Koka requires an empty effect row. Why?

$$\frac{\Gamma \vdash e : \sigma \quad \alpha \notin \text{ftv}(\Gamma)}{\Gamma \vdash e : \forall \alpha. \sigma} \quad \frac{\Gamma \vdash e : \forall \alpha. \sigma}{\Gamma \vdash e : \sigma[\alpha \mapsto T']}$$

Where ftv refers to the free type variables in the context.

B.3.2 Expressions

We freely write τ to mean that a type has an empty effect row. That is, we use τ and a shorthand for $\langle \rangle \tau$. The Δ stands for an arbitrary effect row. We start with everything but the handlers and elaborations and put them in a separate section.

$$\frac{\Gamma_V(x) = \Delta \tau}{\Gamma \vdash x : \Delta \tau} \quad \frac{\Gamma \vdash e : \Delta \tau}{\Gamma \vdash \{e\} : \Delta \tau} \quad \frac{\Gamma \vdash e_1 : \Delta \tau \quad \Gamma_V, x : \tau \vdash e_2 : \Delta \tau'}{\Gamma \vdash \text{let } x = e_1; e_2 : \Delta \tau'}$$

$$\overline{\Gamma \vdash () : \Delta ()} \quad \overline{\Gamma \vdash \text{true} : \Delta \text{Bool}} \quad \overline{\Gamma \vdash \text{false} : \Delta \text{Bool}}$$

$$\frac{\Gamma_V, x_1 : T_1, \dots, x_n : T_n \vdash c : T \quad T_i = \langle \rangle \tau_i}{\Gamma \vdash \text{fn}(x_1 : T_1, \dots, x_n : T_n) T \{e\} : \Delta (T_1, \dots, T_n) \rightarrow T}$$

$$\frac{\Gamma \vdash e_1 : \Delta \text{Bool} \quad \Gamma \vdash e_2 : \Delta \tau \quad \Gamma \vdash e_3 : \Delta \tau}{\Gamma \vdash \text{if } e_1 \{e_2\} \text{ else } \{e_3\} : \Delta \tau}$$

$$\frac{\Gamma \vdash e : (\tau_1, \dots, \tau_n) \rightarrow \Delta \tau \quad \Gamma \vdash e_i : \Delta \tau_i}{\Gamma \vdash e(e_1, \dots, e_n) : \Delta \tau}$$

B.3.3 Declarations and Modules

The modules are gathered into Γ_M and the variables that are in scope are gathered in Γ_V . Each module has the type of its public declarations. Note that these are not accumulative; they only contain the bindings generated by that declaration. Each declaration has the type of both private and public bindings. Without modifier, the public declarations are empty, but with the `pub` keyword, the private bindings are copied into the public declarations.

$$\frac{\Gamma_{i-1} \vdash m_i : \Gamma_{m_i} \quad \Gamma_{M,i} = \Gamma_{M,i-1}, \Gamma_{m_i}}{\Gamma_0 \vdash m_1 \dots m_n : ()}$$

$$\frac{\Gamma_{i-1} \vdash d_i : (\Gamma'_i; \Gamma'_{\text{pub},i}) \quad \Gamma_i = \Gamma_{i-1}, \Gamma'_i \quad \Gamma \vdash \Gamma'_{\text{pub},1}, \dots, \Gamma'_{\text{pub},n}}{\Gamma_0 \vdash \text{mod } x \{d_1 \dots d_n\} : (x : \Gamma)}$$

$$\frac{\Gamma \vdash d : \Gamma'}{\Gamma \vdash d : (\Gamma'; \varepsilon)} \quad \frac{\Gamma \vdash d : \Gamma'}{\Gamma \vdash \text{pub } d : (\Gamma'; \Gamma')} \quad \overline{\Gamma \vdash \text{import } x : \Gamma_M(x)}$$

$$\frac{f_i = \forall \alpha. (\tau_{i,1}, \dots, \tau_{i,n_i}) \rightarrow \alpha x \quad \Gamma'_V = x_1 : f_1, \dots, x_m : f_m}{\Gamma \vdash \text{type } x \{x_1(\tau_{1,1}, \dots, \tau_{1,n_1}), \dots, x_m(\tau_{m,1}, \dots, \tau_{m,n_m})\} : \Gamma'}$$

$$\frac{\Gamma \vdash e : T}{\Gamma \vdash \text{let } x = e : (x : T)}$$

B.3.4 Algebraic Effects and Handlers

Effects are declared with the **effect** keyword. The signatures of the operations are stored in Γ_Φ . The types of the arguments and resumption must all have no effects.

A handler must have operations of the same signatures as one of the effects in the context. The names must match up, as well as the number of arguments and the return type of the expression, given the types of the arguments and the resumption. The handler type then includes the handled effect ϕ , an “input” type τ and an “output” type τ' . In most cases, these will be at least partially generic.

The handle expression will simply add the handled effect to the effect row of the inner expression and use the input and output type.

$$\frac{s_i = op_i(\tau_{i,1}, \dots, \tau_{i,n_i}) : \tau_i \quad \Gamma'_\Phi(x) = \{s_1, \dots, s_n\}}{\Gamma \vdash \mathbf{effect} \ x \ \{s_1, \dots, s_n\} : \Gamma'}$$

$$\frac{\Gamma \vdash e_h : \mathbf{handler} \ \phi \ \tau \ \tau' \quad \Gamma \vdash e_c : \langle \phi | \Delta \rangle \ \tau}{\Gamma \vdash \mathbf{handle} \ e_h \ e_c : \Delta \ \tau'}$$

$$\frac{\begin{array}{l} A(\phi) \quad \Gamma_\Phi(\phi) = \{s_1, \dots, s_n\} \quad \Gamma, x : \tau \vdash e_{\text{ret}} : \tau' \\ \left[\begin{array}{l} s_i = x_i(\tau_{i,1}, \dots, \tau_{i,m_i}) \rightarrow \tau_i \quad o_i = x_i(x_{i,1}, \dots, x_{i,m_i}) \{e_i\} \\ \Gamma_V, \text{resume} : (\tau_i) \rightarrow \tau', x_{i,1} : \tau_{i,1}, \dots, x_{i,i_m} : \tau_{i,i_m} \vdash e_i : \tau' \end{array} \right]_{1 \leq i \leq n} \end{array}}{\Gamma \vdash \mathbf{handler} \ \{\text{return}(x)\{e_{\text{ret}}\}, o_1, \dots, o_n\} : \mathbf{handler} \ \phi \ \tau \ \tau'}$$

B.3.5 Higher-Order Effects and Elaborations

The declaration of higher-order effects is similar to first-order effects, but with exclamation marks after the effect name and all operations. This will help distinguish them from first-order effects.

Elaborations are of course similar to handlers, but we explicitly state the higher-order effect $x!$ they elaborate and which first-order effects Δ they elaborate into. The operations do not get a continuation, so the type checking is a bit different there. As arguments, they take the effectless types they specified along with the effect row Δ . Elaborations are not added to the value context, but to a special elaboration context mapping the effect identifier to the row of effects to elaborate into.

The **elab** expression then checks that an elaboration for all higher-order effects in the inner expression are in scope and that all effects they elaborate into are handled.

$$\frac{s_i = op_i!(\tau_{i,1}, \dots, \tau_{i,n_i}) : \tau_i \quad \Gamma'_\Phi(x!) = \{s_1, \dots, s_n\}}{\Gamma \vdash \mathbf{effect} \ x! \ \{s_1, \dots, s_n\} : \Gamma'}$$

$$\frac{\begin{array}{l} \Gamma_\Phi(x!) = \{s_1, \dots, s_n\} \quad \Gamma'_E(x!) = \Delta \\ \left[\begin{array}{l} s_i = x_i!(\tau_{i,1}, \dots, \tau_{i,m_i}) \ \tau_i \quad o_i = x_i!(x_{i,1}, \dots, x_{i,m_i}) \{e_i\} \\ \Gamma, x_{i,1} : \Delta \ \tau_{i,1}, \dots, x_{i,n_i} : \Delta \ \tau_{i,n_i} \vdash e_i : \Delta \ \tau_i \end{array} \right]_{1 \leq i \leq n} \end{array}}{\Gamma \vdash \mathbf{elaboration} \ x! \ \rightarrow \Delta \ \{o_1, \dots, o_n\} : \Gamma'}$$

INFO: Later, we could add more precise syntax for which effects need to be present in the arguments of the elaboration operations.

INFO: It is not possible to elaborate only some higher-order effects. We could change the behaviour to allow this later.

$$\frac{[\Gamma_E(\phi) \subseteq \Delta]_{\phi \in H(\Delta')} \quad \Gamma \vdash e : \Delta' \tau \quad \Delta = A(\Delta')}{\Gamma \vdash \text{elab } e : \Delta \tau}$$

B.4 Desugaring

To simplify the reduction rules, we simplify the AST by desugaring some constructs. This transform is given by a fold over the syntax tree with the following operation:

$$\begin{aligned} D(\text{fn}(x_1 : T_1, \dots, x_n : T_n) T \{e\}) &= \lambda x_1, \dots, x_n. e \\ D(\text{let } x = e_1; e_2) &= (\lambda x. e_2)(e_1) \\ D(e_1; e_2) &= (\lambda _. e_2)(e_1) \\ D(\{e\}) &= e \end{aligned}$$

B.5 Semantics

The semantics of Elaine are defined as reduction semantics.

We use two separate contexts to evaluate expressions. The E context is for all constructs except effect operations, such as **if**, **let** and function applications. The X_{op} context is the context in which a handler can reduce an operation op .

$$\begin{aligned} E ::= & [] \mid E(e_1, \dots, e_n) \mid v(v_1, \dots, v_n, E, e_1, \dots, e_m) \\ & \mid \text{if } E \{e\} \text{ else } \{e\} \\ & \mid \text{let } x = E; e \mid E; e \\ & \mid \text{handle}[E] e \mid \text{handle}[v] E \\ & \mid \text{elab}[E] e \mid \text{elab}[v] E \end{aligned}$$

$$\begin{aligned} X_{op} ::= & [] \mid X_{op}(e_1, \dots, e_n) \mid v(v_1, \dots, v_n, X_{op}, e_1, \dots, e_m) \\ & \mid \text{if } X_{op} \{e_1\} \text{ else } \{e_2\} \\ & \mid \text{let } x = X_{op}; e \mid X_{op}; e \\ & \mid \text{handle}[X_{op}] e \mid \text{handle}[h] X_{op} \text{ if } op \notin h \\ & \mid \text{elab}[X_{op}] e \mid \text{elab}[\epsilon] X_{op} \text{ if } op! \notin e \end{aligned}$$

TODO: Add some explanation

$$\begin{aligned} c(v_1, \dots, v_n) &\longrightarrow \delta(c, v_1, \dots, v_n) \\ &\quad \text{if } \delta(c, v_1, \dots, v_n) \text{ defined} \\ (\lambda x_1, \dots, x_n. e)(v_1, \dots, v_n) &\longrightarrow e[x_1 \mapsto v_1, \dots, x_n \mapsto v_n] \\ \text{if true } \{e_1\} \text{ else } \{e_2\} &\longrightarrow e_1 \\ \text{if false } \{e_1\} \text{ else } \{e_2\} &\longrightarrow e_2 \\ \text{handle}[h] v &\longrightarrow e[x \mapsto v] \end{aligned}$$

$$\begin{aligned}
& \text{where } \text{return}(x)\{e\} \in H \\
\text{handle}[h] X_{op}[op(v_1, \dots, v_n)] & \longrightarrow e[x_1 \mapsto v_1, \dots, x_n \mapsto v_n, \text{resume} \mapsto k] \\
& \text{where } op(x_1, \dots, x_n)\{e\} \in h \\
& k = \lambda y. \text{handle}[h] X_{op}[y] \\
\text{elab}[\epsilon] v & \longrightarrow v \\
\text{elab}[\epsilon] X_{op!}[op!(e_1, \dots, e_n)] & \longrightarrow \text{elab}[\epsilon] X_{op!}[e[x_1 \mapsto e_1, \dots, x_n \mapsto e_n]] \\
& \text{where } op!(x_1, \dots, x_n)\{e\} \in \epsilon
\end{aligned}$$

B.6 Standard Library

Elaine does not include any operators. This choice was made to simplify parsing of the language. For the lack of operators, any manipulation of primitives needs to be done via the standard library of built-in functions.

These functions reside in the `std` module, which can be imported like any other module with the `use` statement to bring its contents into scope.

The full list of functions available in the `std` module, along with their signatures and descriptions, is given in Figure B.2.

	Name	Type signature		Description
Arithmetic	add	fn (Int, Int)	Int	addition
	sub	fn (Int, Int)	Int	subtraction
	neg	fn (Int)	Int	negation
	mul	fn (Int, Int)	Int	multiplication
	div	fn (Int, Int)	Int	division
	modulo	fn (Int, Int)	Int	modulo
	pow	fn (Int, Int)	Int	exponentiation
Comparisons	eq	fn (Int, Int)	Bool	equality
	neq	fn (Int, Int)	Bool	inequality
	gt	fn (Int, Int)	Bool	greater than
	geq	fn (Int, Int)	Bool	greater than or equal
	lt	fn (Int, Int)	Bool	less than
	leq	fn (Int, Int)	Bool	less than or equal
Boolean operations	not	fn (Bool)	Bool	boolean negation
	and	fn (Bool, Bool)	Bool	boolean and
	or	fn (Bool, Bool)	Bool	boolean or
String operations	concat	fn (Bool, Bool)	Bool	string concatenation
Conversions	show_int	fn (Int)	String	integer to string
	show_bool	fn (Bool)	String	integer to string

Figure B.2: Overview of the functions in the `std` module in Elaine.